

Deployment to Kubernetes in Bukalapak

Kubernetes and Cloud Native Bandung Meetup #3
Lo.Ka.Si Coffee & Space

16 November 2019

BukaLapak



Muhammad Saiful Islam

Cloud Engineer @ Bukalapak, 2019 — now

Site Reliability Engineer @ Bukalapak, 2018 — 2019

Software Engineering Intern @ Bukalapak, 2017

S.Tr.Kom. from Politeknik Negeri Bandung, 2014 — 2018



Introduction to Kubernetes

Giri Kuncoro, 13 September 2019



Migration Journey to Kubernetes

Didiet Agus Pambudiono, 24 October 2019

both are accessible from

<https://github.com/cloudnative-id/kubernetes-cloud-native-bandung>

A quick **recap**...

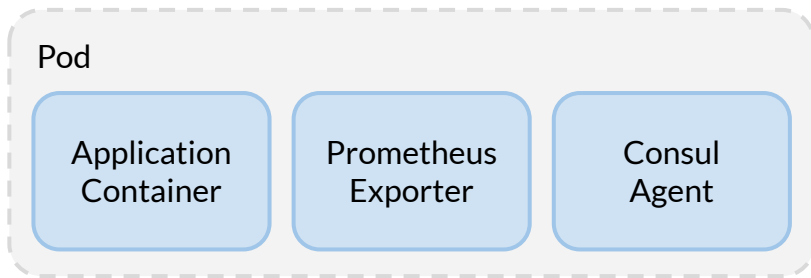
Kubernetes

Pod

Small groups of
containers and **volumes**

Tight coupling —
live and die together

Shared **networking**
between containers



Kubernetes Deployment

Declarative updates for pods

Specify a **desired state**,
controller handles the rest

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: lulu-api
spec:
  replicas: 2
  template:
    metadata:
      labels:
        project: lulu
        service: api
    spec:
      containers:
        - name: lulu-api
          ports:
            - containerPort: 8080
          resources:
            limits:
              memory: 2Gi
              cpu: 2
            requests:
              memory: 0.5Gi
              cpu: 1
```

Kubernetes Service

Exposes a group of pods into
a **network service**


Provides **load balancing** and
a **DNS name** to be accessed
by other workloads

```
apiVersion: v1
kind: Service
metadata:
  name: lulu-api
spec:
  selector:
    project: lulu
    service: api
  ports:
    - protocol: TCP
      port: 8080
  type: ClusterIP
```

Horizontal Pod Autoscaler

Automatically **add (or remove) pods**
when needed

Based on **CPU/memory metrics**
with **custom metrics support**



>70

Ruby services

>150

Go services

>150

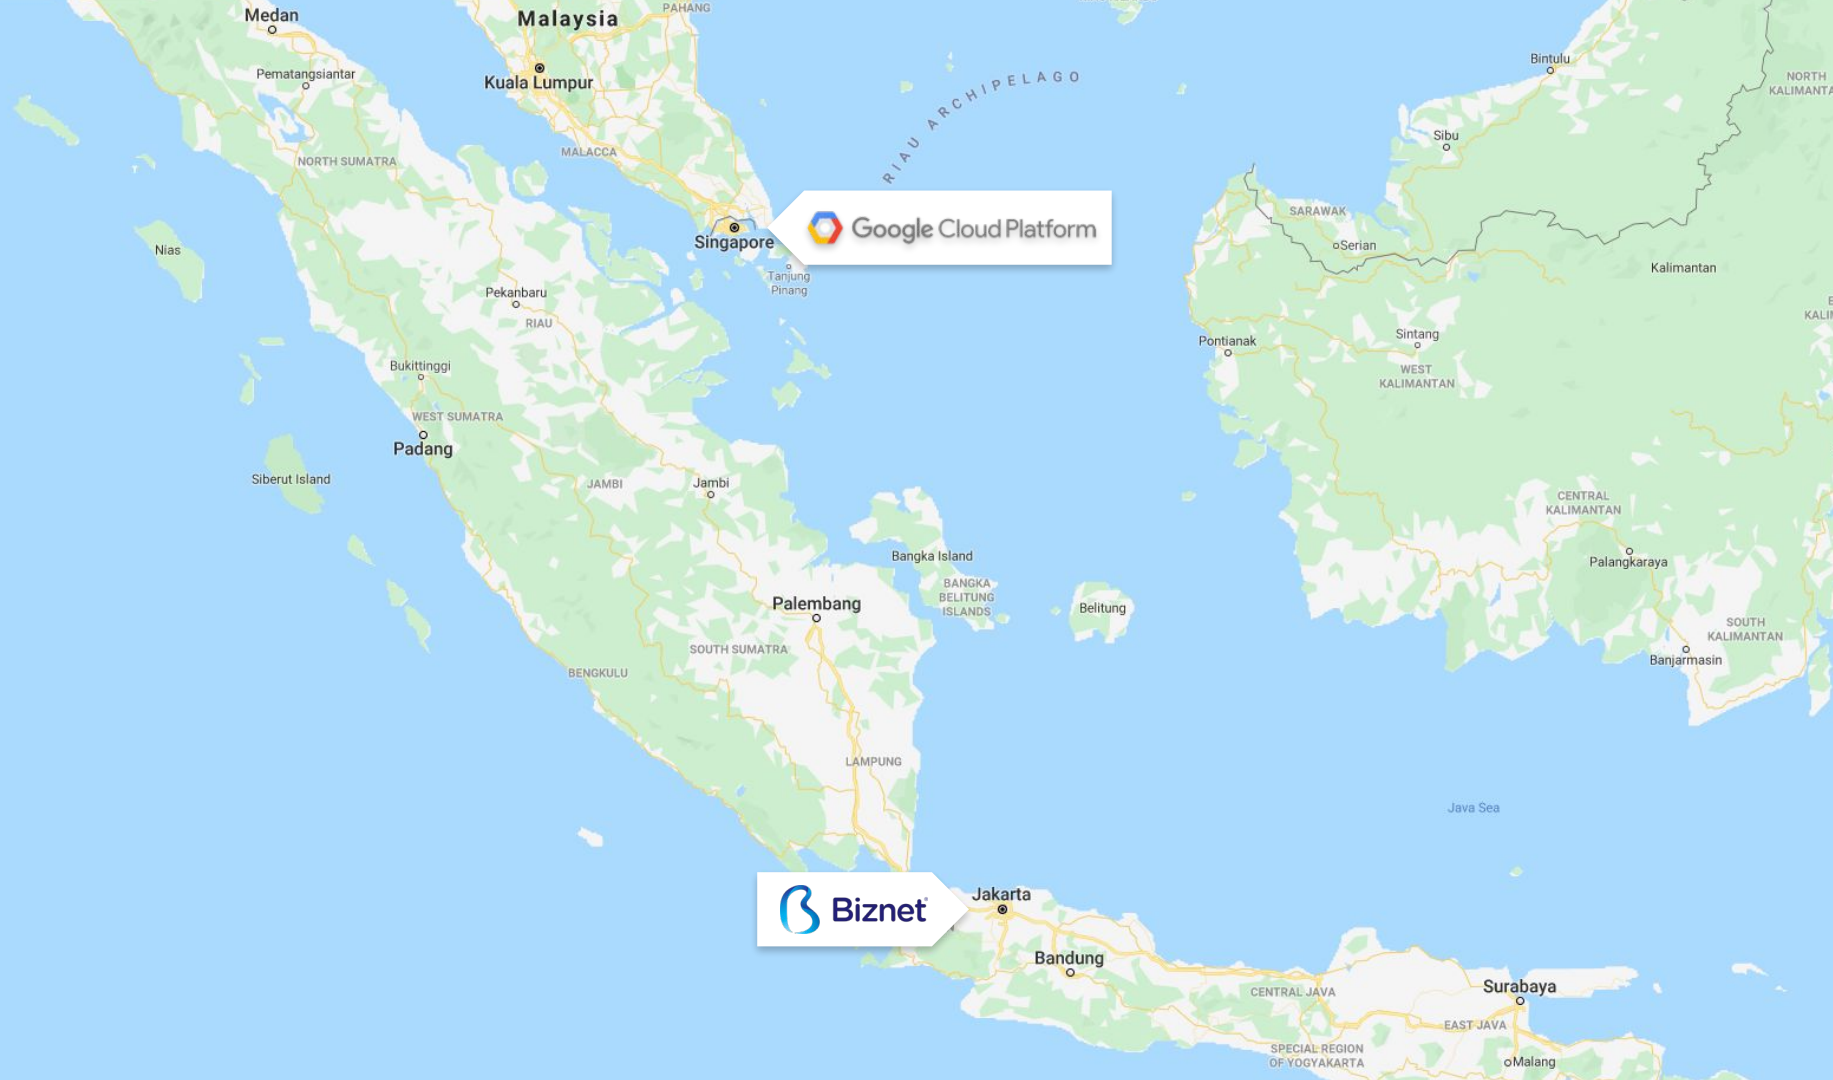
Java/NodeJS/Python/Elixir/
other services

>1.700

deployments

>20.000

pods running



Google Cloud Platform


Biznet

So, how do we **deploy**
our microservices to our clusters?

Test

Pre-production

Production



stages:

- test
- before-build-preproduction
- build-preproduction
- before-deploy-preproduction
- deploy-preproduction
- after-deploy-preproduction
- before-build-canary
- build-canary
- before-deploy-canary
- deploy-canary
- before-deploy-production
- deploy-production

test:

stage: test

image: \$CI_REGISTRY/bukalapak/gudang/golang:2.4.1

script:

- make test
- GIT_BRANCH=\$CI_COMMIT_REF_NAME goveralls -coverprofile=coverage.out

-endpoint=\$COVERALLS_ENDPOINT

lint:

stage: test


image: \$CI_REGISTRY/bukalapak/gudang/golang:2.4.1

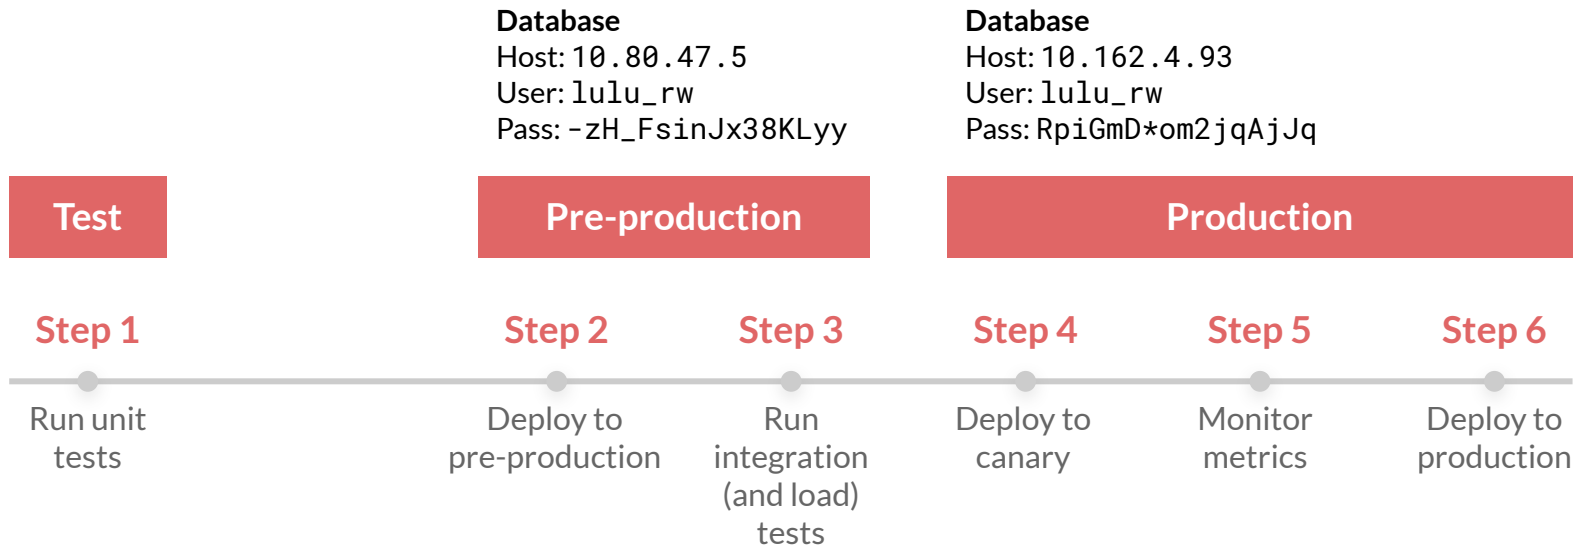
script:

- curl -sfL <https://install.goreleaser.com/github.com/golangci/golangci-lint.sh> | sh -s --
- b \$(go env GOPATH)/bin v1.16.0
- golangci-lint --version
- golangci-lint run



```
compile-preproduction:
  stage: before-build-preproduction
  image: $CI_REGISTRY/bukalapak/gudang/golang:2.4.1
  only:
    refs:
      - master
      - /^release.+/
  script:
    - make compile
  artifacts:
    paths:
      - deploy/_output
```

```
build-image-preproduction:
  stage: build-preproduction
  image: $CI_REGISTRY/bukalapak/gudang/docker:1.0.0
  only:
    refs:
      - master
      - /^release.+/
  script:
    - export VERSION=${CI_COMMIT_TAG:-$CI_COMMIT_SHA}-preproduction
    - docker login -u "$CI_REGISTRY_USER" -p "$CI_REGISTRY_PASSWORD" "$CI_REGISTRY"
    - make build-image push
```





Should we build it
for each environment?



<https://12factor.net/config>

```
// Copyright 2017, Google, Inc.  
// Licensed under the Apache License, Version 2.0 (the "License")  
var http = require('http');  
var server = http.createServer(function (request, response) {  
  const language = 'English';  
  const API_KEY = '123-456-789';  
  response.write(`Language: ${language}\n`);  
  response.write(`API Key: ${API_KEY}\n`);  
  response.end(`\n`);  
});  
server.listen(3000);
```

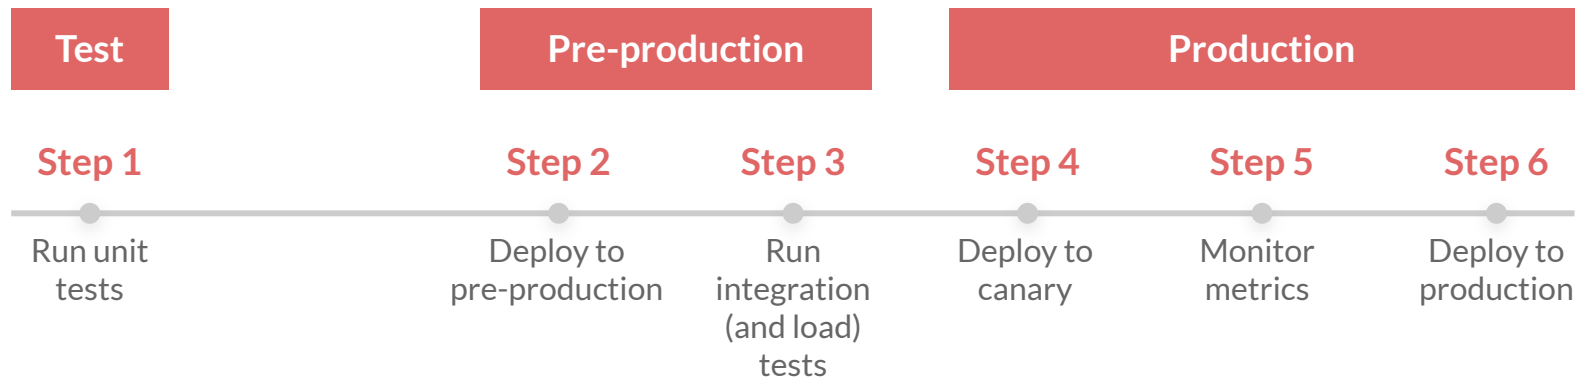
<https://medium.com/google-cloud/kubernetes-configmaps-and-secrets-68d061f7ab5b>

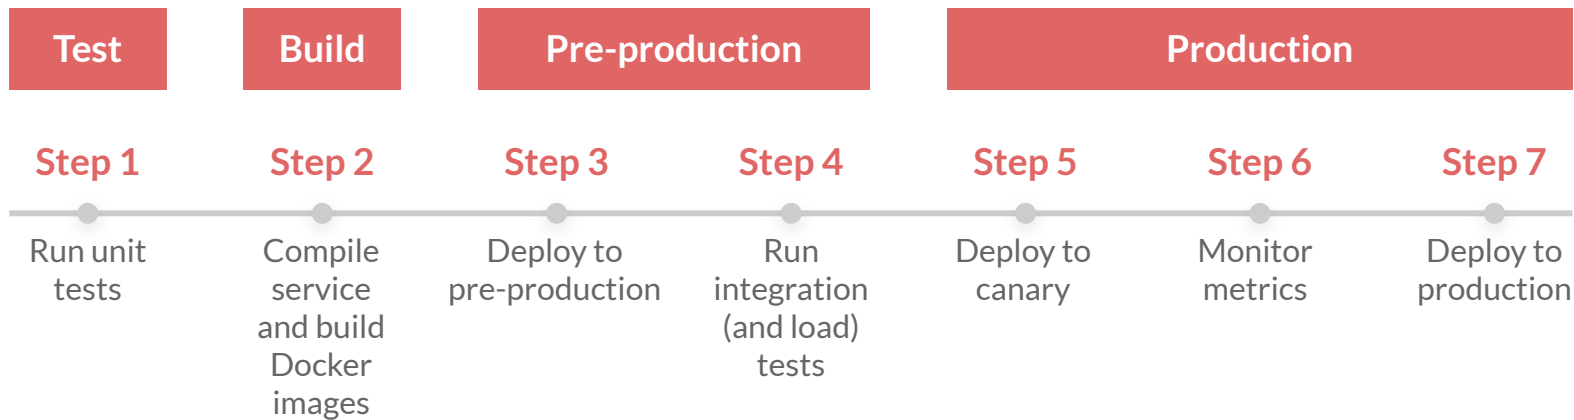
```
// Copyright 2017, Google, Inc.  
// Licensed under the Apache License, Version 2.0 (the "License")  
var http = require('http');  
var server = http.createServer(function (request, response) {  
  const language = 'English';  
  const API_KEY = '123-456-789';  
  response.write(`Language: ${language}\n`);  
  response.write(`API Key: ${API_KEY}\n`);  
  response.end(`\n`);  
});  
server.listen(3000);
```

<https://medium.com/google-cloud/kubernetes-configmaps-and-secrets-68d061f7ab5b>

```
// Copyright 2017, Google, Inc.  
// Licensed under the Apache License, Version 2.0 (the "License")  
var http = require('http');  
var server = http.createServer(function (request, response) {  
  const language = process.env.LANGUAGE;  
  const API_KEY = process.env.API_KEY;  
  response.write(`Language: ${language}\n`);  
  response.write(`API Key: ${API_KEY}\n`);  
  response.end(`\n`);  
});  
server.listen(3000);
```

<https://medium.com/google-cloud/kubernetes-configmaps-and-secrets-68d061f7ab5b>





stages:

- test
- before-build-preproduction
- build-preproduction
- before-deploy-preproduction
- deploy-preproduction
- after-deploy-preproduction
- before-build-canary
- build-canary
- before-deploy-canary
- deploy-canary
- before-deploy-production
- deploy-production

stages:

- test
- ~~— before-build-preproduction~~
- ~~— build-preproduction~~
- before-deploy-preproduction
- deploy-preproduction
- after-deploy-preproduction
- ~~— before-build-canary~~
- ~~— build-canary~~
- before-deploy-canary
- deploy-canary
- before-deploy-production
- deploy-production

stages:

- test
- before-build
- build
- before-deploy-preproduction
- deploy-preproduction
- after-deploy-preproduction
- before-deploy-canary
- deploy-canary
- before-deploy-production
- deploy-production

```
compile-preproduction:
  stage: before-build-preproduction
  image: $CI_REGISTRY/bukalapak/gudang/golang:2.4.1
  only:
    refs:
      - master
      - /^release.+/
  script:
    - make compile
  artifacts:
    paths:
      - deploy/_output
```

```
build-image-preproduction:
  stage: build-preproduction
  image: $CI_REGISTRY/bukalapak/gudang/docker:1.0.0
  only:
    refs:
      - master
      - /^release.+/
  script:
    - export VERSION=${CI_COMMIT_TAG:-$CI_COMMIT_SHA}-preproduction
    - docker login -u "$CI_REGISTRY_USER" -p "$CI_REGISTRY_PASSWORD" "$CI_REGISTRY"
    - make build-image push
```

```
compile:
  stage: before-build
  image: $CI_REGISTRY/bukalapak/gudang/golang:2.4.1
  only:
    refs:
      - master
      - /^release.+/
  script:
    - make compile
  artifacts:
    paths:
      - deploy/_output

build-image:
  stage: build
  image: $CI_REGISTRY/bukalapak/gudang/docker:1.0.0
  only:
    refs:
      - master
      - /^release.+/
  script:
    - export VERSION=${CI_COMMIT_TAG:-$CI_COMMIT_SHA}
    - docker login -u "$CI_REGISTRY_USER" -p "$CI_REGISTRY_PASSWORD" "$CI_REGISTRY"
    - make build-image push
```

How do we **deploy and link**
the configuration **at runtime**?

Kubernetes ConfigMap

Bind **environment variables**
to Pods' containers
at runtime

Separate **configuration**
from **code**

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: lulu-config
data:
  DATABASE_HOST: 10.80.47.5
  DATABASE_USERNAME: lulu_rw
  DATABASE_PASSWORD: -zH_FsinJx38KLyy
```


Kubernetes ConfigMap

Bind **environment variables**
to Pods' containers
at runtime

Separate **configuration**
from **code**

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: lulu-config
data:
  DATABASE_HOST: 10.80.47.5
  DATABASE_USERNAME: lulu_rw
  DATABASE_PASSWORD: zH_FsinJx38KLyy
```

Kubernetes Secret

Store **sensitive data**
in clusters

Give control over
how sensitive data is used

Reduce the risk
of exposing the data
to **unauthorized users**

```
apiVersion: v1
kind: Secret
metadata:
  name: lulu-secret
data:
  DATABASE_PASSWORD: LXpIX0ZzaW5KeDM4S0x5eQ==
```

Link them all


Store **sensitive data**
in clusters

Give control over
how sensitive data is used


Reduce the risk
of exposing the data
to **unauthorized users**

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: lulu-api
spec:
  template:
    spec:
      containers:
        - name: lulu-api
          envFrom:
            - configMapRef:
                name: lulu-config
            - secretRef:
                name: lulu-secret
```

How do we **manage**
the configuration?



<https://12factor.net/config>



Minerva

Stores service **configuration files** and **manifest templates**

Ensures **separation** between code and configuration

Provides CI/CD **pipeline and tooling**

minerva└─ **services**└─ **lulu**└─ **config**

└─ kubeconfig.yaml

└─ kubelize

└─ **preproduction**

└─ web

└─ **production**

└─ web

└─ **manifest**└─ **web**

└─ deployment.yaml

└─ hpa.yaml

└─ service.yaml

└─ resourcequota.yaml

└─ ingress.yaml


```
apiVersion: v1
clusters:
- name: preproduction
  cluster:
    certificate-authority-data: <certificate-for-your-cluster>
    server: https://<cluster-endpoint>
- name: production
  cluster:
    certificate-authority-data: <certificate-for-your-cluster>
    server: https://<cluster-endpoint>
```

```
projectName: sample-examples
services:
- name: web
  port: "8888"
  language: go
  environments:
  ...
  minCpu: "1"
  maxCpu: "2"
  minMem: "0.5"
  maxMem: "2"
  averageUtilization: 65
  rollingUpdate:
    maxUnavailable: 10%
    maxSurge: 25%
  healthCheck:
    delay: "10"
    healthz: healthz
    ready: healthz
    period: 4
    timeout: "2"
    liveFailThreshold: 1
    liveSuccThreshold: 1
    readFailThreshold: 1
    readSuccThreshold: 2
```

```
environments:
- name: canary
  replica: "1"
- name: production
  replica: "5"
  minReplicas: "5"
  maxReplicas: "20"
  quotaRequestsCpu: 25
  quotaLimitsCpu: 50
  quotaRequestsMem: 12.5
  quotaLimitsMem: 50
- name: preproduction
  replica: "1"
  minReplicas: "1"
  maxReplicas: "5"
  quotaRequestsCpu: 6.25
  quotaLimitsCpu: 12.5
  quotaRequestsMem: 3.125
  quotaLimitsMem: 12.5
```

ENV=preproduction

DATABASE_HOST=10.80.47.5

DATABASE_USERNAME=lulu_rw

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: {{.config.projectName}}-{{.service.name}}-{{.environment.name}}
  namespace: {{.config.projectName}}
  ...
spec:
  replicas: {{.environment.replica}}
  ...
template:
  metadata: ...
  spec:
    ...
    containers:
    - name: {{.config.projectName}}-{{.service.name}}
      image: {{.variable.CI_REGISTRY}}/bukalapak/{{.config.projectName}}/{{.service.name}}:{{.variable.VERSION}}
      ...
    envFrom:
      - configMapRef:
          name: {{.config.projectName}}-{{.service.name}}-{{or .variable.SUFFIX .environment.name}}
          optional: false
      - secretRef:
          name: {{.config.projectName}}-{{.service.name}}-{{or .variable.SUFFIX .environment.name}}
          optional: false
```




Hashicorp Vault

Stores **arbitrary key/value secrets**

Encrypts the secrets prior to writing them to persistent storage

Controls **access** to tokens, passwords, ...

<https://www.vaultproject.io>



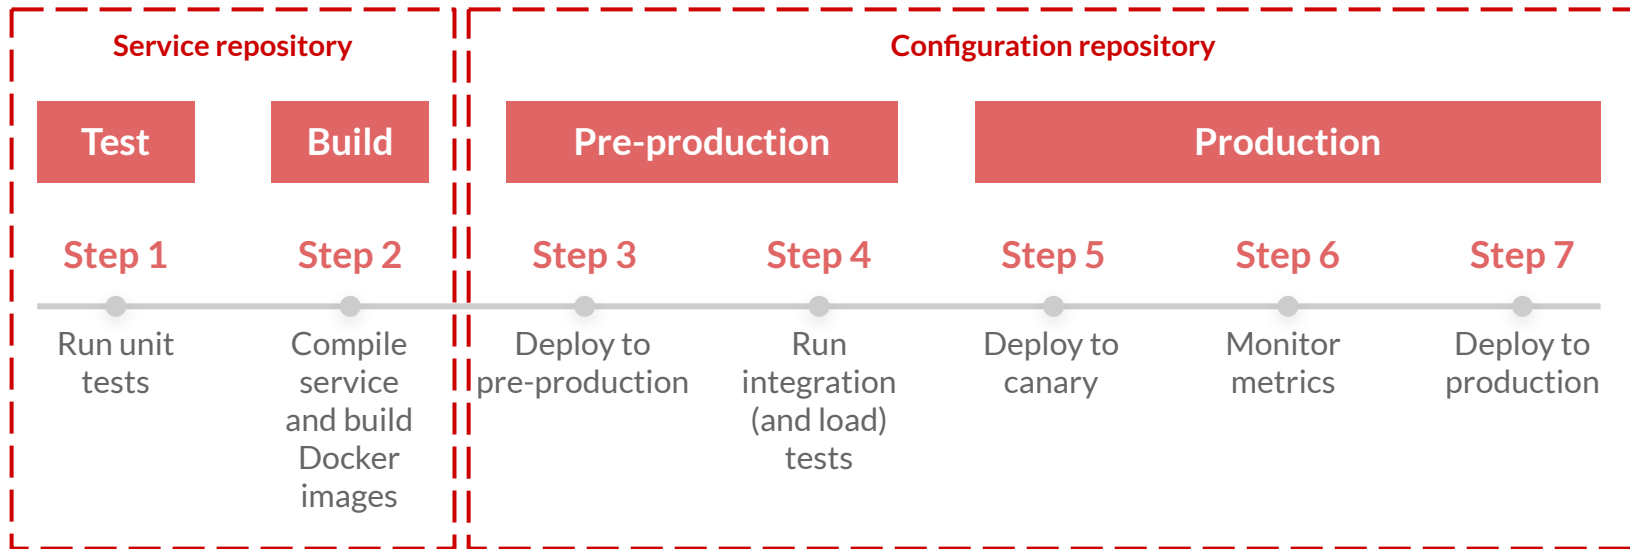
```
include:  
- project: 'bukalapak/minerva'  
  ref: v1.0.0  
  file: 'shared/pipeline-template/.gitlab-ci-general.yaml'
```

```
test: ...
```

```
lint: ...
```

```
compile: ...
```

```
build: ...
```




```
preproduction-manifest:
  extends: .generate-and-validate-manifest
  stage: before-preproduction-deployment
  variables:
    ENV: preproduction

preproduction-configuration:
  extends: .configuration-deployment
  stage: before-preproduction-deployment
  variables:
    ENV: preproduction
```

```
deploy-preproduction:
  extends: .deploy-template
  stage: preproduction-deployment
  variables:
    ENV: preproduction
    FILEPATH: deploy/_output/$ENV/manifest
  environment:
    name: preproduction
  artifacts:
    paths:
      - ${FILEPATH}/.rollout_check.list

preproduction-done:
  extends: .releaselog-done
  stage: after-preproduction-deployment
  variables:
    NOTIFICATION_ENV: preproduction
```

Then ... we need to **rollback**.

```
compile:
  stage: before-build
  image: $CI_REGISTRY/bukalapak/gudang/golang:2.4.1
  only:
  refs:
  - master
  - /^release.+/
  script:
  - make compile
  artifacts:
    expire_in: 1 day
    paths:
    - deploy/_output
```

```
build-image:
  stage: build
  image: $CI_REGISTRY/bukalapak/gudang/docker:1.0.0
  only:
  refs:
  - master
  - /^release.+/
  script:
  - export VERSION=${CI_COMMIT_TAG:-$CI_COMMIT_SHA}
  - docker login -u "$CI_REGISTRY_USER" -p "$CI_REGISTRY_PASSWORD" "$CI_REGISTRY"
  - make build-image push
```

How to link between
code and config **version**?

peru


Includes code from another **repository**

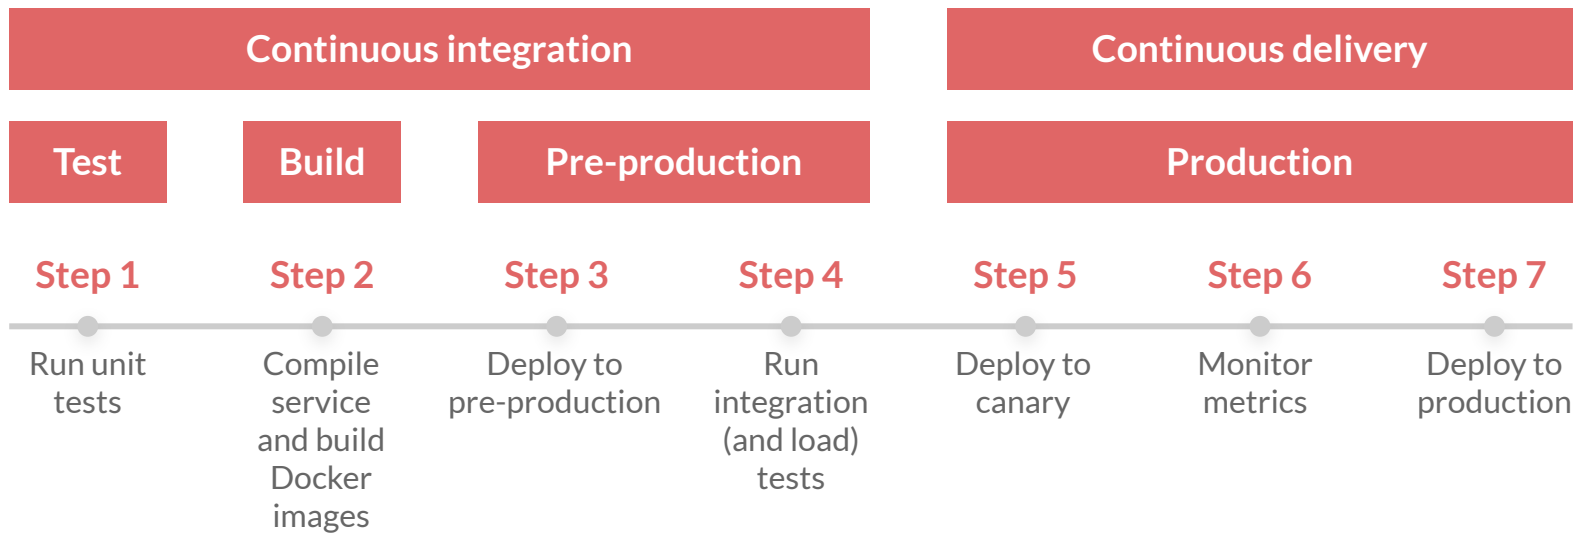
Helps track **exact version** of dependencies

so the history is always **reproducible**

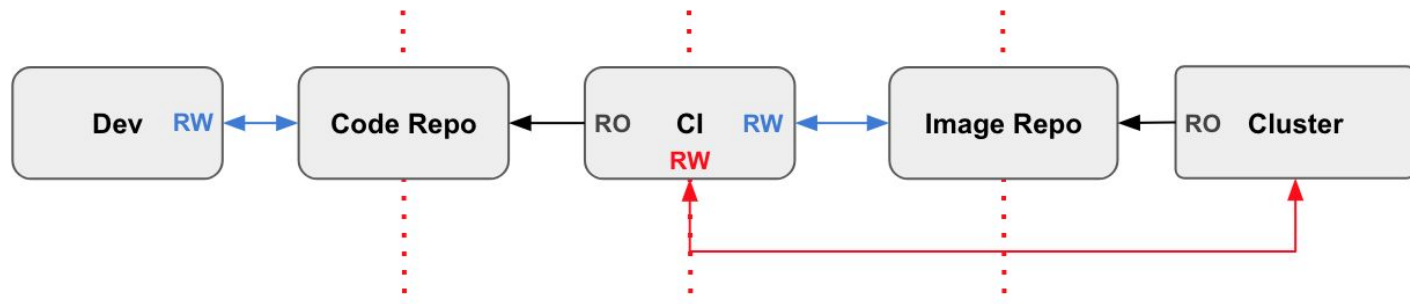
<https://github.com/buildinspace/peru>

```
imports:  
  minerva: minerva/  
  
git module minerva:  
  url: https://gitlab.bl.internal/bukalapak/minerva.git  
  rev: 2f9f7ef5aaa60d55279ffd255d42556f30e3507a
```

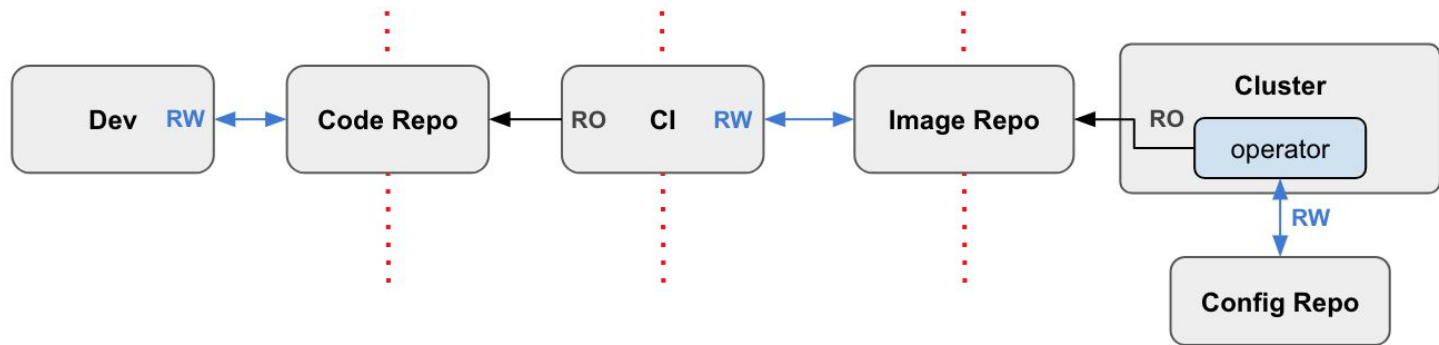




Is it ... **GitOps** already?



<https://www.weave.works/technologies/gitops/>



<https://www.weave.works/technologies/gitops/>

Want to improve that?

Scratch your own itch; let's work together ;)

Bukalapak

Thank you!

Muhammad Saiful Islam

@saifulwebid on the Internet

saiful.islam@bukalapak.com

Bukalapak