




Running Istio in TVLK Data Production

Imre Nagi

@imrenagi

Software Engineer @TVLK DATA

About Me?



Imre Nagi

@imrenagi

- Software Engineer @traveloka data
- ex - Software Engineer Intern @eBay & @CERN
- Indonesia Docker Community Leader
- One of Jakarta Kubernetes Organizer team member

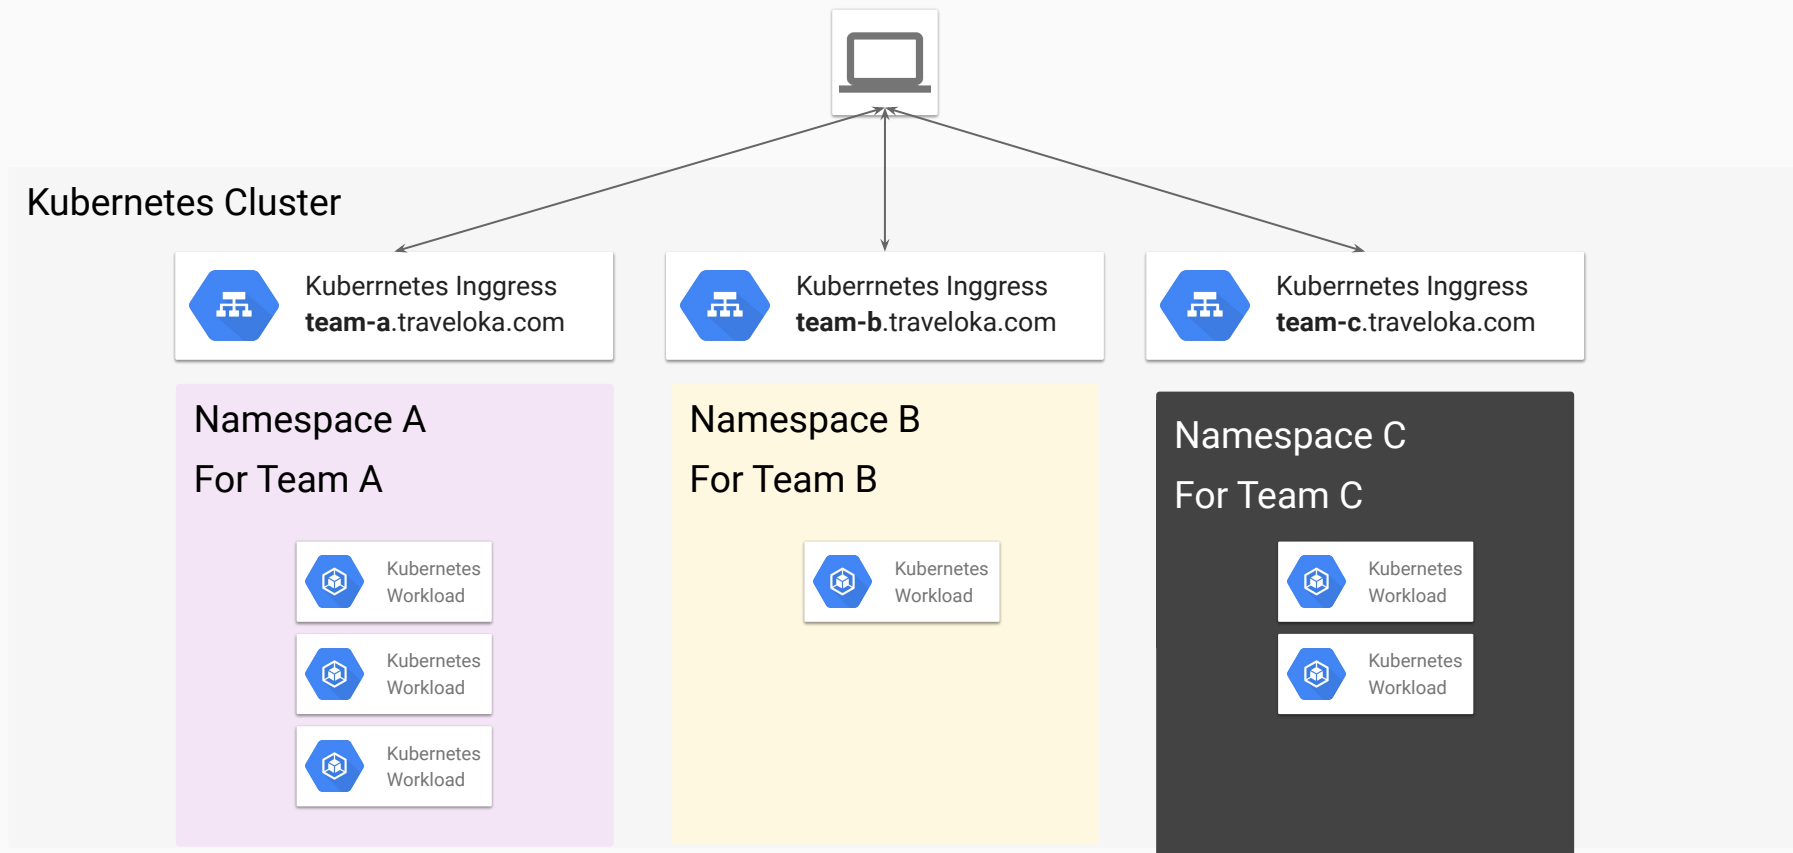
Agenda

- Kubernetes In Traveloka
- Istio
- Running Istio @traveloka
- Lesson learned

Kubernetes in TVLK

Kubernetes, Long story short

- About 2 years in Data Team Production
- Multiple teams in TVLK Data are using Kubernetes to run multiple type of loads.
 - API,
 - Visualization Dashboard,
 - GPU and Machine Learning
- At least, handles tens of thousands of request per second
- Cost Optimization (up to 60%)
 - Node and pod autoscaling
 - Better CPU utilizations




* Similar configuration between prod & staging cluster






“Everything is just working fine until one team decided to start decoupling their applications into multiple smaller services.”



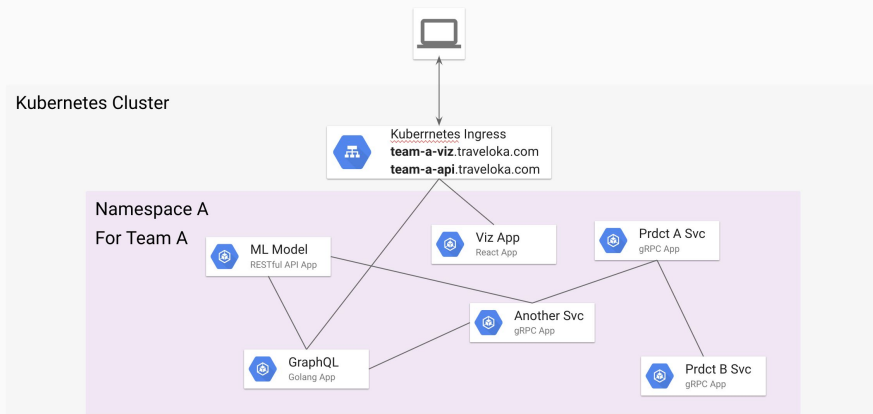
Kubernetes Cluster



Namespace A For Team A



Managing Service Mesh is Freakin' Complicated



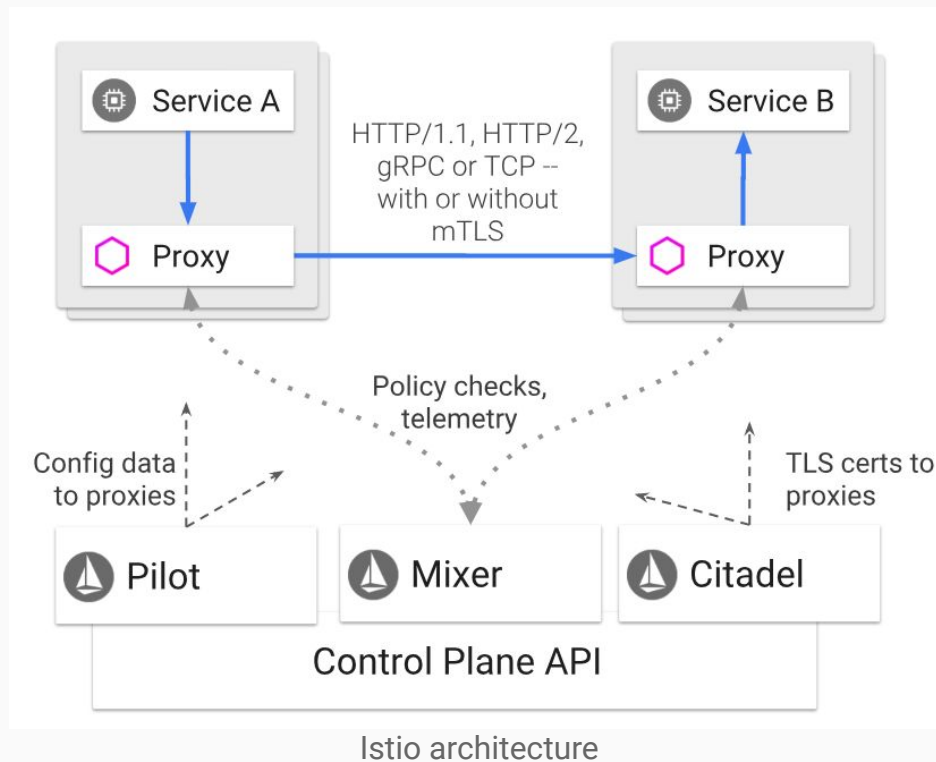
- Securing inter-service communication
- Unstandardized monitoring approach
 - Different monitoring tools used by different programming lang
 - Need better way to monitor RPC/API call between services (RED metrics, and Distributed Tracing)
- gRPC Load balancing doesn't work out of the box..

Istio to simplify service mesh in
microservice

Why istio?

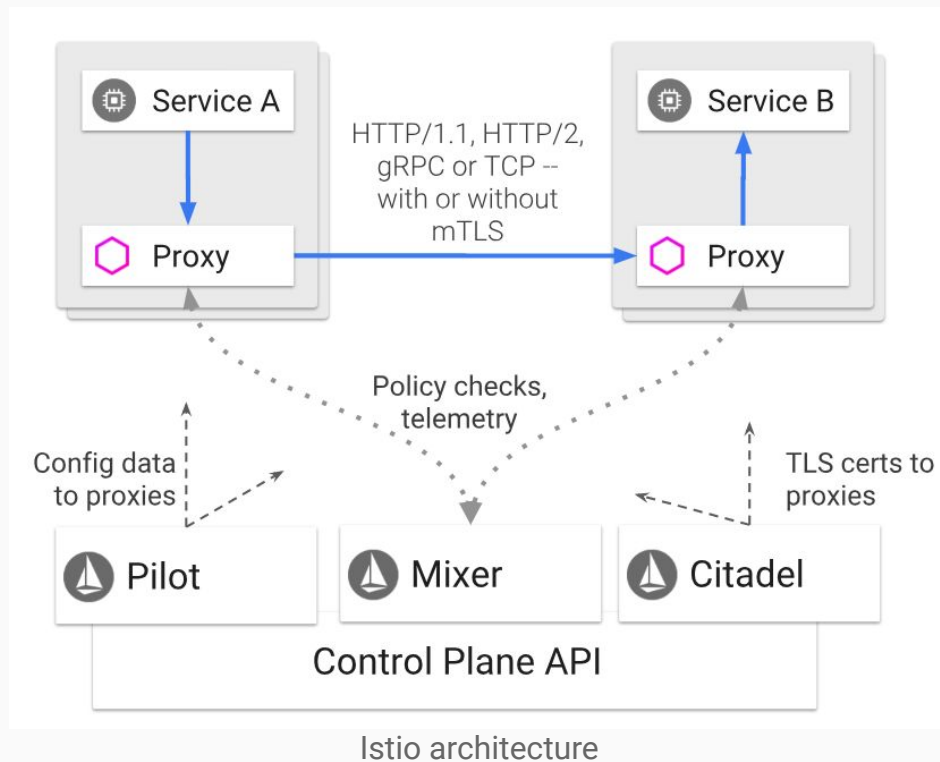
Istio features

- Traffic Management
- Security
- Monitoring

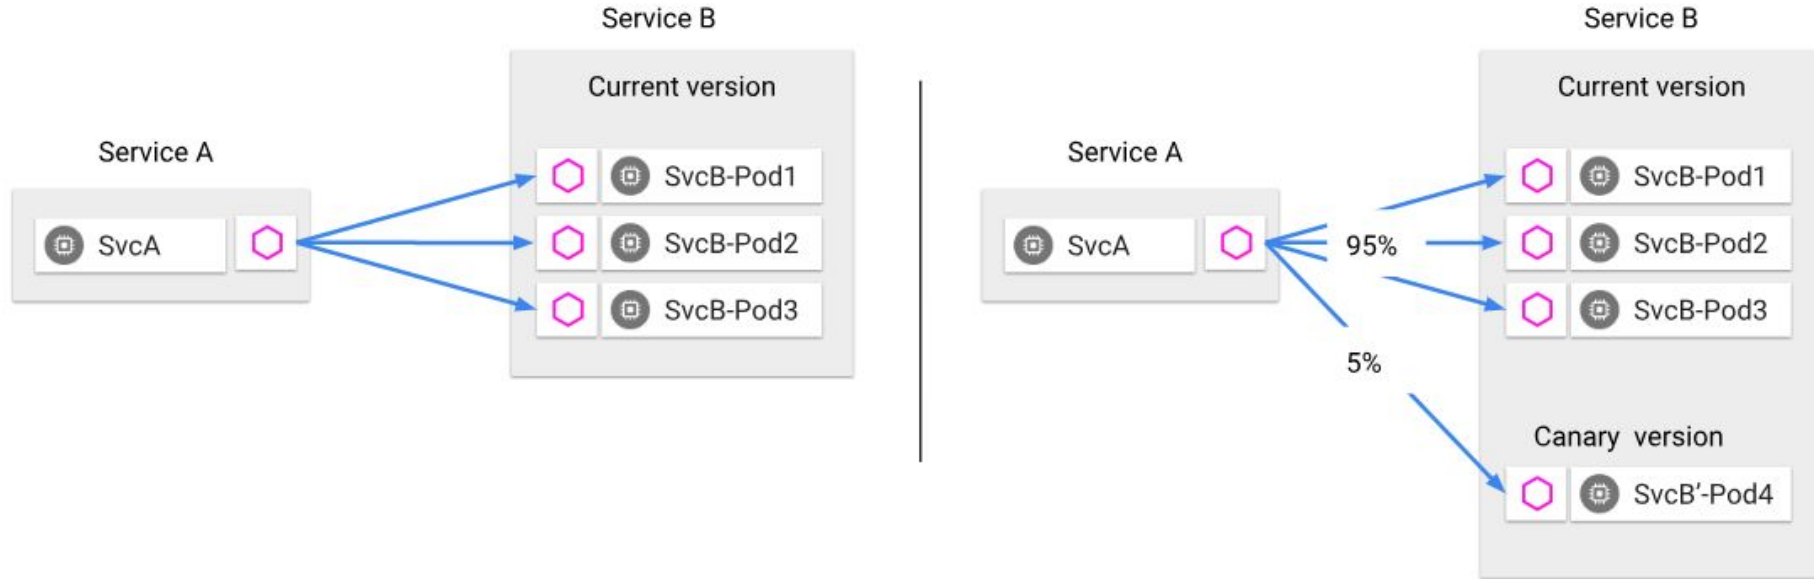


Istio Traffic Management

- Intelligent routing
 - A/B tests
 - Canary deployments
- Resiliency
 - timeouts,
 - retries,
 - circuit breakers,
 - etc.

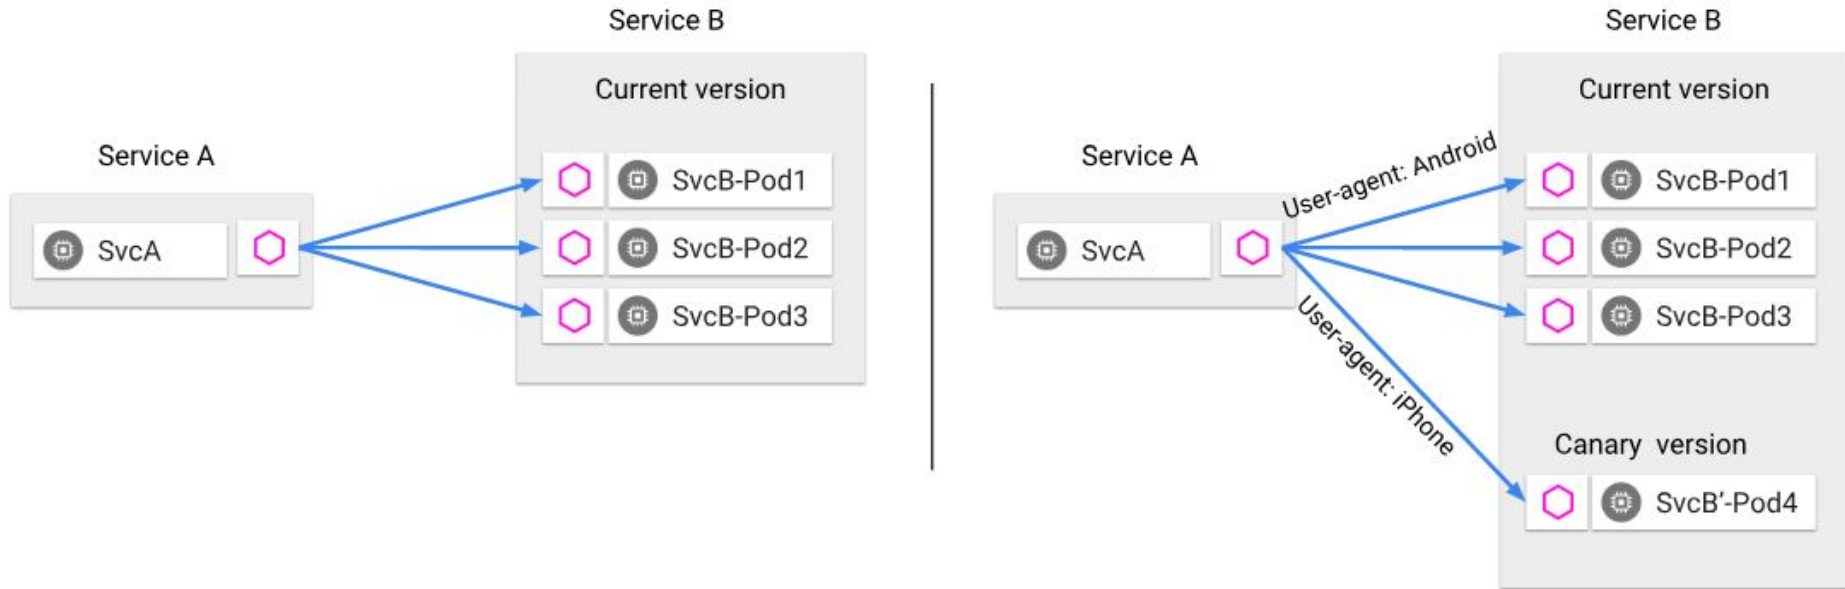


Splitting Traffic Based on Traffic Proportion



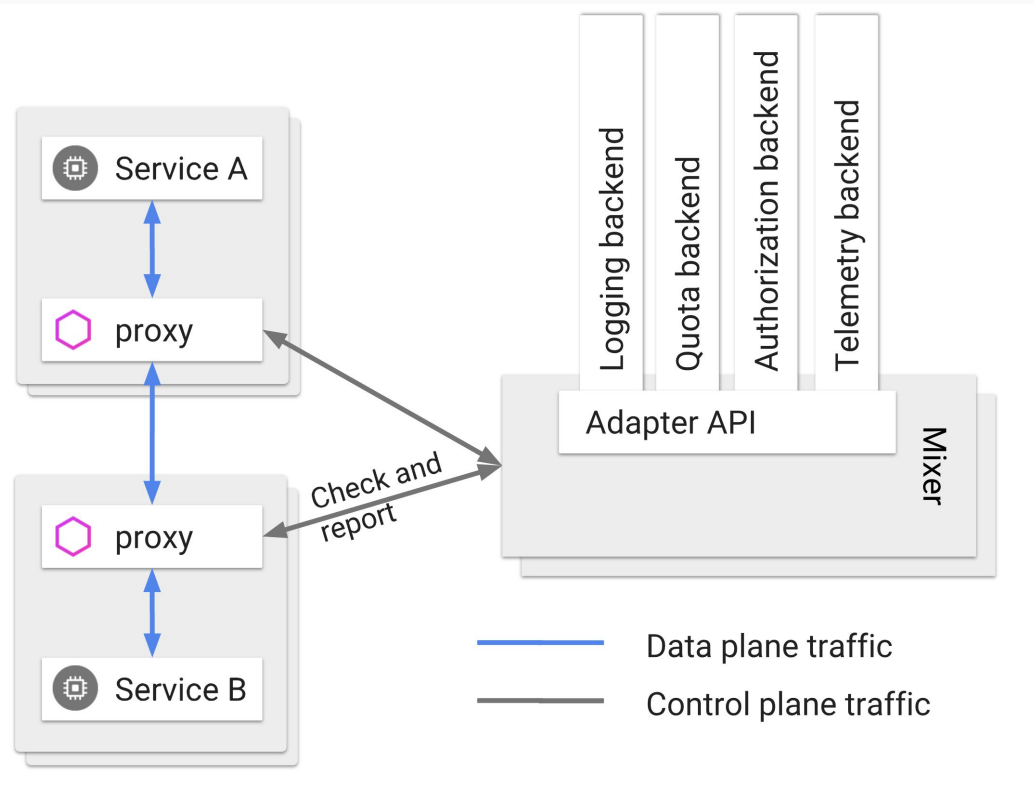
Traffic splitting decoupled from infrastructure scaling - proportion of traffic routed to a version is independent of number of instances supporting the version

Splitting Traffic Based on It's Request Content



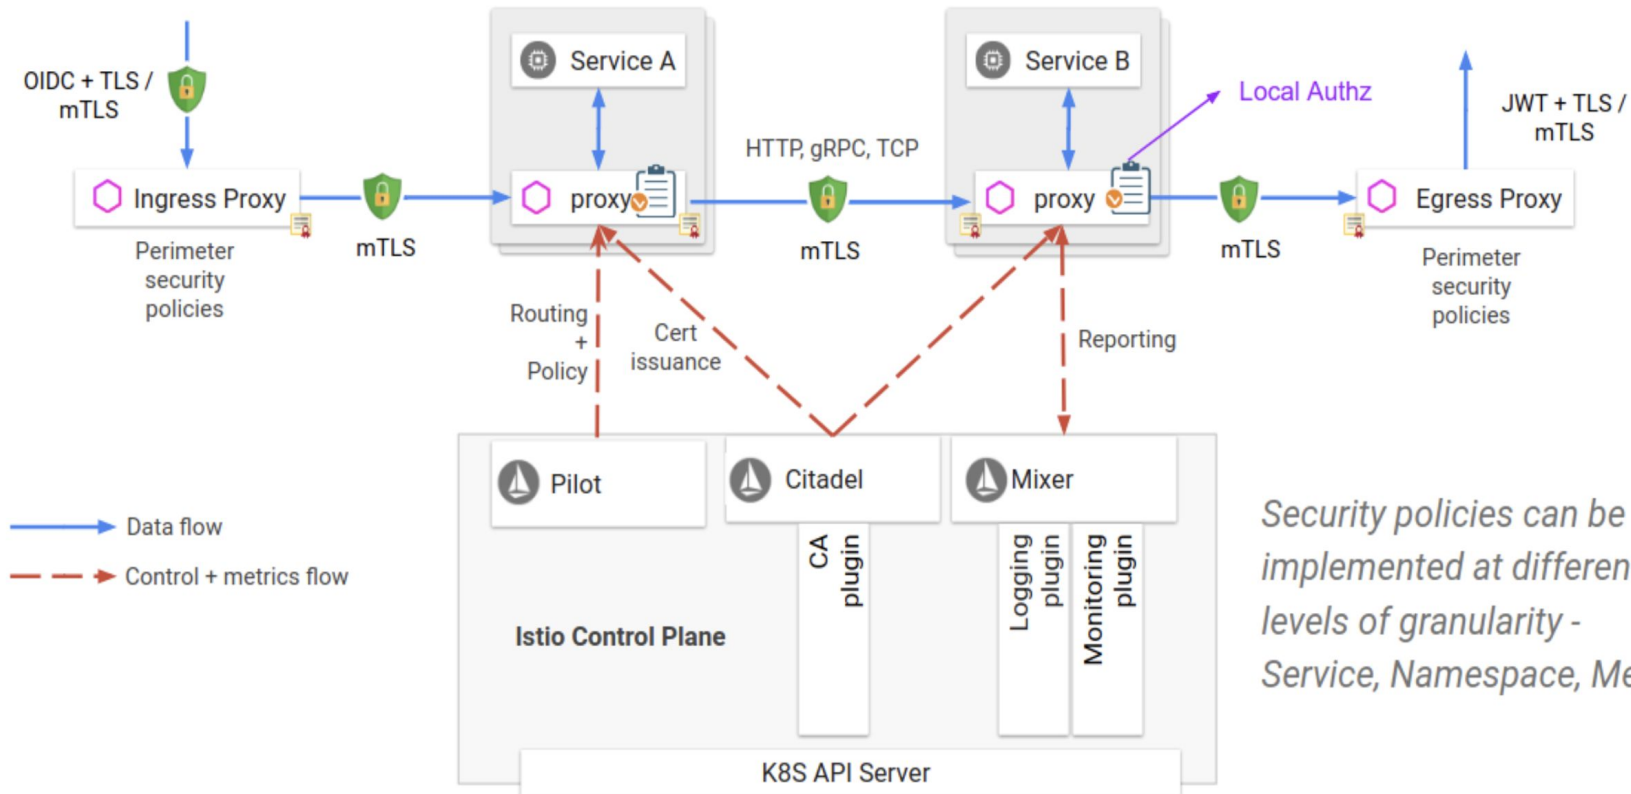
Content-based traffic steering - The content of a request can be used to determine the destination of a request

Istio Mixer for Collecting Metrics



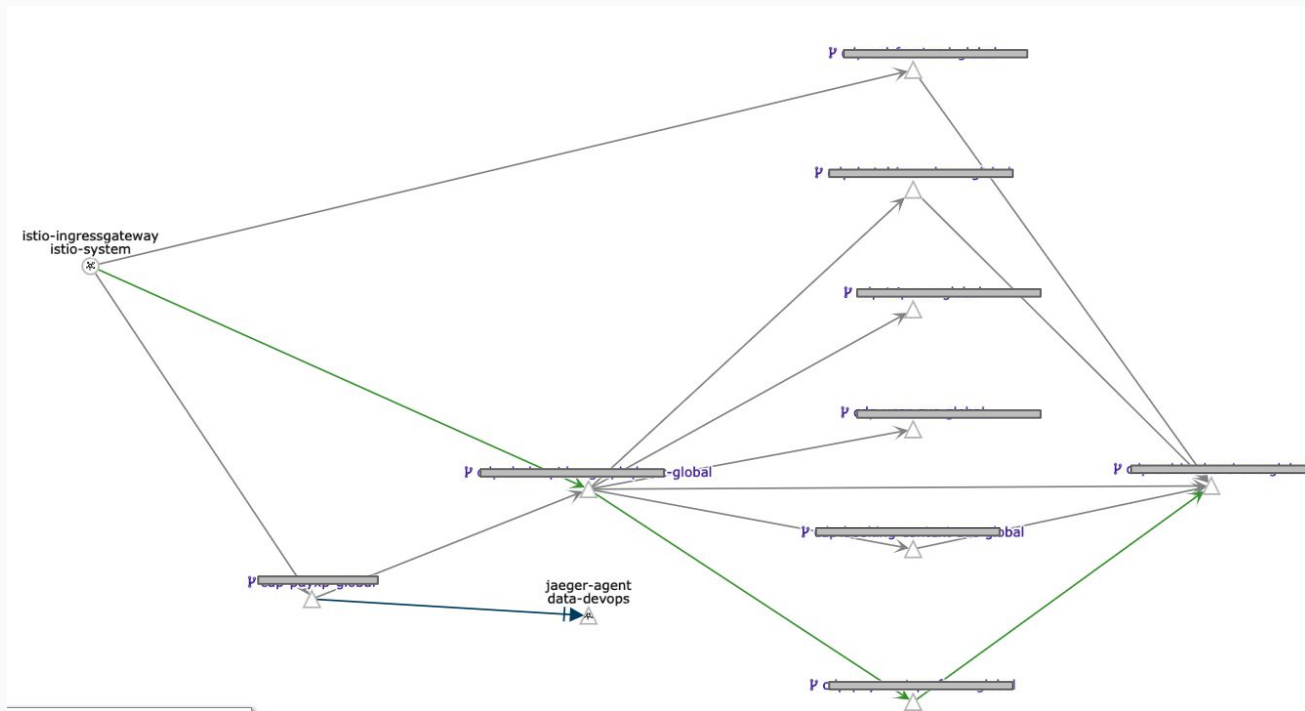
- Mixer is platform independent component
- Enforces access control and usage policies across the mesh.
- Collect telemetry data from envoy
- No vendor lock-in. This enables istio to interface with variety of backend

Istio Citadel for Cert Issuance



Running Istio @ Traveloka Data

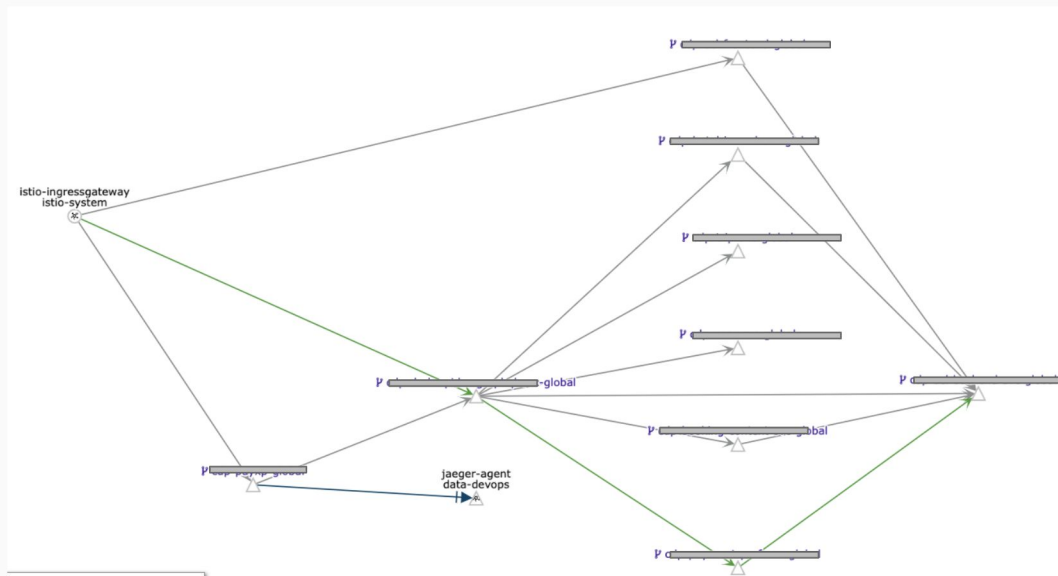
App that is running in istio?



- Mostly Golang + gRPC
- Some RESTful API
- GraphQL server
- Some visualizations & dashboards

Issues in Monitoring the Mesh (Before Istio)

- Unstandardized monitoring metrics.
 - Each service owner has his/her own approach in collecting the metrics with Datadog (latency, request count, etc)
- Hard to separate metrics produced by different version (canary and stable)
- Hard to see the client side metrics





Monitoring Dashboard (With Istio)

- RED* (**R**ate, **E**rror, **D**uration) metrics are collected by Istio Mixer and stored in Prometheus
- Those metrics are visualized in Grafana
 - Supported query by regex
 - Plug and Play some reusable dashboards





*<https://www.weave.works/blog/the-red-method-key-metrics-for-microservices-architecture/>

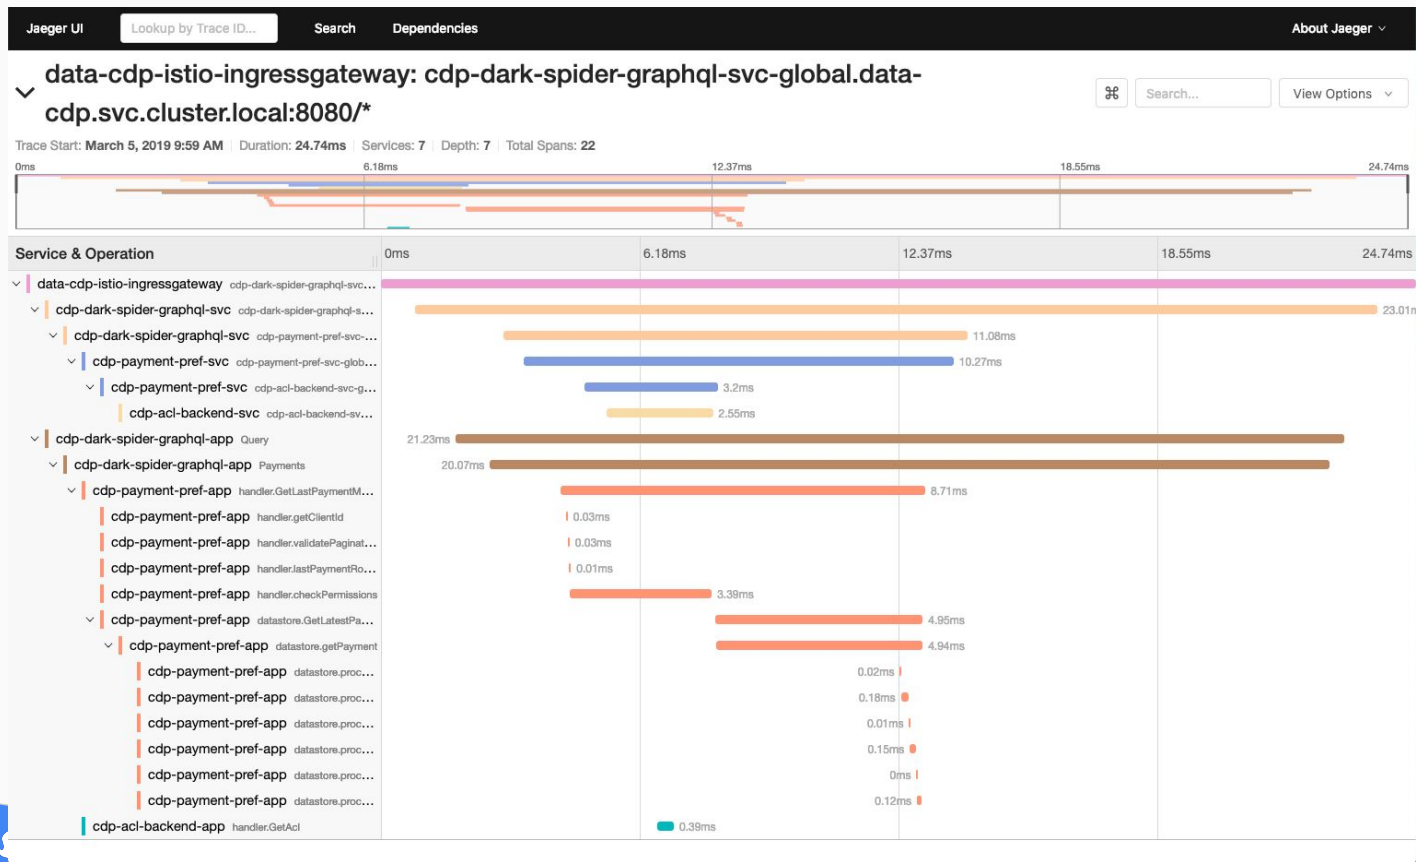
Stable & Canary monitoring



Client side metrics



Distributed Tracing with Jaeger



Kubernetes Cluster



Istio Gateway Controller
For Team A



Istio Gateway Controller
For Team B

Namespace A For Team A



Kubernetes
Workload



Kubernetes
Workload



Kubernetes
Workload

Namespace B For Team B



Kubernetes
Workload



Kubernetes
Workload

Dedicated Istio Gateway for Each Team

Benefits:

- Decentralized
- Failure isolation
- Cost Isolation

DevOps Helm Chart Template

- Expose relevant istio's configuration related to the traffic management, security policy, etc to service developer
- Reduced developer effort to maintain all of their istio and kubernetes manifests, Increased productivity.
- Split into two helm-chart
 - Application helm chart: app/pod deployment
 - Infra helm chart: infra, service routing, and security setting

default values supplied for templates/* files

```
app:
  istioProxy:
    enabled: true
  name: example-app
  track: stable
  env: stg
  baseImageName: gcr.io/[REDACTED]/java-maven-app
  imageTag: c1448e9
  containerPort: 8080
# see <protocol> supported here: https://istio.io/docs/setup/kubernetes/spec-requirements/
containerPortProtocol: http
```

define volumes in the pod

```
volumes:
- name: example-app-secret-volume
  secret:
    secretName: example-app-secrets
    # set optional true if the secret is not that important
    optional: true
```

```
# mount specific volume to app's container
volumeMounts:
- name: example-app-secret-volume
  readOnly: true
  mountPath: /etc/example-app/secrets

# the app's container resource spec
resources:
  requests:
    cpu: "0.5"
    memory: "1G"
  limits:
    cpu: "1"
    memory: "2G"

# autoscale configuration
autoscale:
  enabled: true
  maxReplica: 10
  minReplica: 1
  targetCPUUtilizationPercentage: 60
```

Expose Service Through Istio Ingress Gateway

This configuration below is normally needed to expose your service to public.

```
infra:
  ingressGateway:
    enabled: true
    httpsRedirect: true
    selector:
      app: istio-ingressgateway # use other selector if you want to have a different gateway
    hosts:
      - "foo.tvlk-data.com"
```

Protect Service Access Within The Same Cluster with mTLS (Mutual TLS)

Normally, you only use TLS Certs managed by Istio Citadel, the Istio component that manage (e.g: renew and deploy) certificates automatically for your services.

```
infra:
# this trafficPolicy is a configuration that applied to <app-name>-global service communication.
# it tells the client which uses istio-proxy, to use this trafficPolicy when communicating with our <app-n
# see https://istio.io/docs/reference/config/istio.networking.v1alpha3/#TrafficPolicy to check the availab
trafficPolicy:
  tls:
    # see: https://istio.io/docs/reference/config/istio.networking.v1alpha3/#TLSSETTINGS-TLSmode
    # `mode` possible values: DISABLE, SIMPLE, MUTUAL, ISTIO_MUTUAL
    mode: ISTIO_MUTUAL

# servicePolicy: is authentication policy enforcement when communicating with the `target` service in our
# technically speaking, the `servicePolicy` field is a `spec` field in istio's Policy object, see: https://
servicePolicy:
  # targets: a list of target service in a shortname.
  targets:
  - name: [redacted]-svc-global
  peers:
  - mtls:
      # mtls set to STRICT in production
      mode: PERMISSIVE # possible values are PERMISSIVE, STRICT
```

Split Traffic to Different Track of Service

Given there are 2 track of services: `stable` and `canary` have been deployed, you can split the traffic based on weight (with total of 100 percent). Here's how you can do that

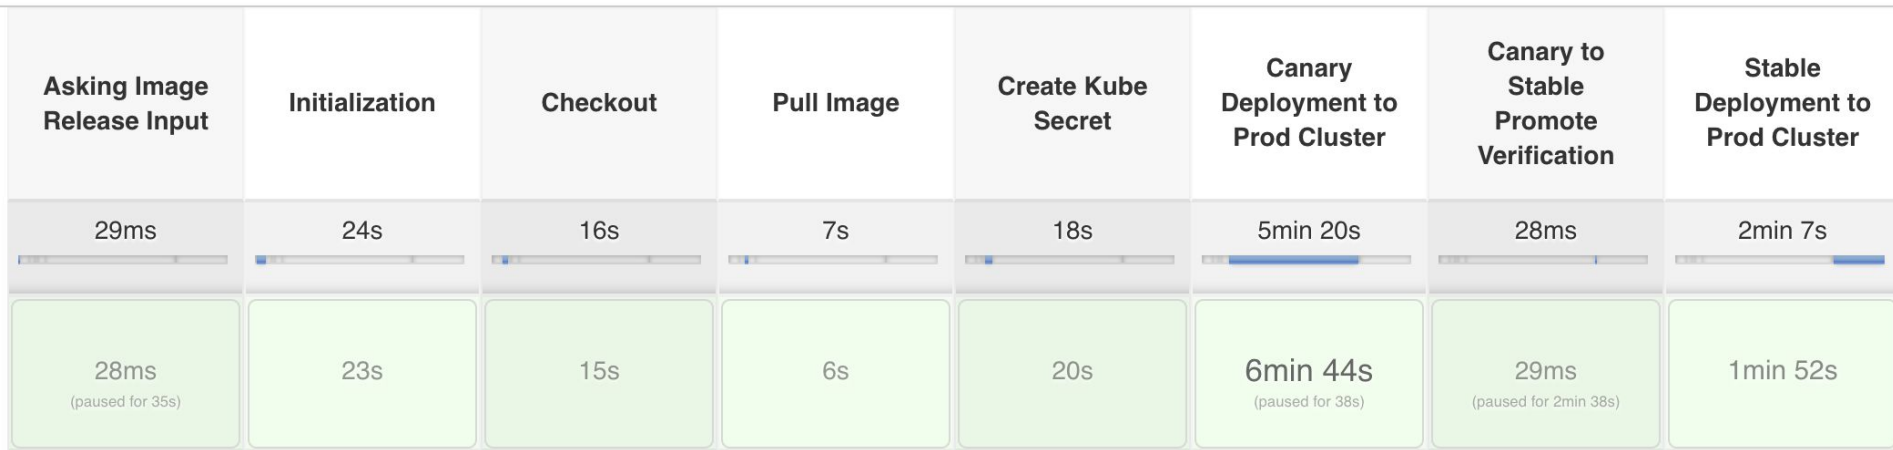
```
infra:
  trafficRoute:
    weight:
      stable: 95
      canary: 5
    canaryCondition:
      match:
        - headers:
            end-user:
              exact: apratama
          key:
            exact: agung
```

Traffic Mirroring/Shadowing to Different Version/Track of Service

Given there are 2 track of services: `stable` and `canary` have been deployed, you can mirror the current `stable` traffic to the `canary` as well, without impact end-user. The response from `canary` service isn't delivered to the upstream (end-user), so you can think of it as fire-and-forget. The only response delivered to the upstream is from `stable`. Here's how to do that:

```
infra:
  trafficRoute:
    weight:
      stable: 100
      canary: 0
    mirror:
      enabled: true
      track: canary
```

Continuous Integration & Delivery



- Decide which image to deploy
- Ask how much traffic goes to canary?
- Verify the canary work as expected?
- Prod deployment

Deploy w/ Makefile

```
stage("Stable Deployment to Prod Cluster") {
    echo "Deploy ${appImage} as stable deployment to cluster"

    sh "make deploy-app TRACK=stable ${commonDeployParams} ARGS=\"--set app.autoscale.enabled=true --timeout ${deployWaitInSeconds} --force\""

    echo "Setup Infra around application"
    sh """
    make deploy-infra ENV=${environment} NAMESPACE=${appNameNamespace} \
    ARGS=\"--set infra.trafficRoute.weight.canary=0 --set infra.trafficRoute.weight.stable=100 --timeout ${deployWaitInSeconds} --force\"
    """

    echo "Cleaning up canary deployment"
    sh """
    sleep 10
    make delete APP_NAME=${appName} TYPE=canary
    """
}
```



Lesson Learned


Lesson learned

- FYI, we are running OSS Istio in GKE.
- At least, our PoC can handle 800 RPS, enough for now
- Some Issues we found
 - Different architecture used in our staging and prod made incident post mortem become harder when we got issue in one of those env. (Hard to reproduce the issue)
 - gRPC mirroring doesn't work as expected in istio.
 - Authentication for human user interaction doesn't exist in Istio.
- Always run enough load test to:
 - Fail faster
 - Tune the autoscaling config
 - Tune monitoring resiliency



Question?

23:18 ngegas APP @Imre ngegas:



@Imre ngegas:



send email to: imre.nagi@traveloka.com for referral 😊

