

Networking Model and Services

In Kubernetes

Arkan from  Bahamify



Jakarta Kubernetes Meetup
11 Mei 2019

Intro

- Early Stage Startup(≤ 1 year development).
4 month MVP. 2 engineer
- Startup School Awardee, YCombinator
- 99% Container Tech, Microservice,
Orchestrated and Built on top of k8s




Overview of K8S Model

Center of K8s: Networking



- using cluster concept
- manage communication and regulation
 - ◆ between apps/container inside cluster
 - ◆ between cluster and external traffic



Objective Menu

- Pod Ip Model
- Services
- Cluster to External Communication
- Network Policy




Pod Ip Model

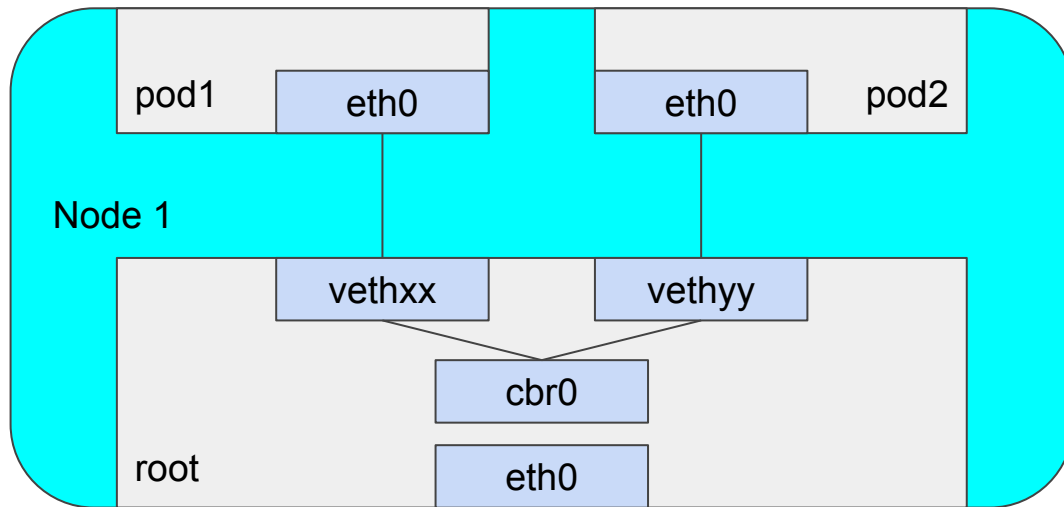
Every Pod Has Real IP Address

Basic Rule

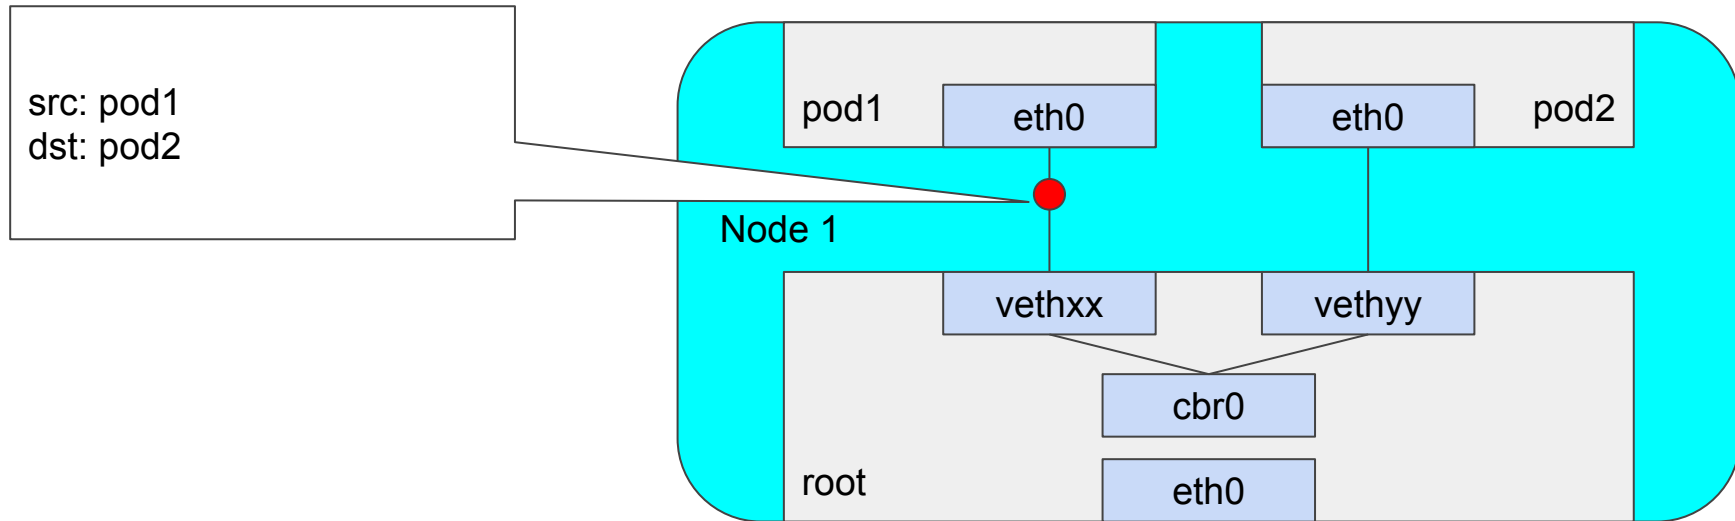
- Pod can be accessed from other pod, no matter where they are (can be distributed across the node, **flat**)
- Every node has CIDR(IP Block)
- K8s **DOESNT CARE** how (L2, L3, overlay, carrier pigeons,dll)
- Pods vs Docker
 - no port mapping
 - no shared machine private IP




Network Namespace




Pod to Pod - Intra Node



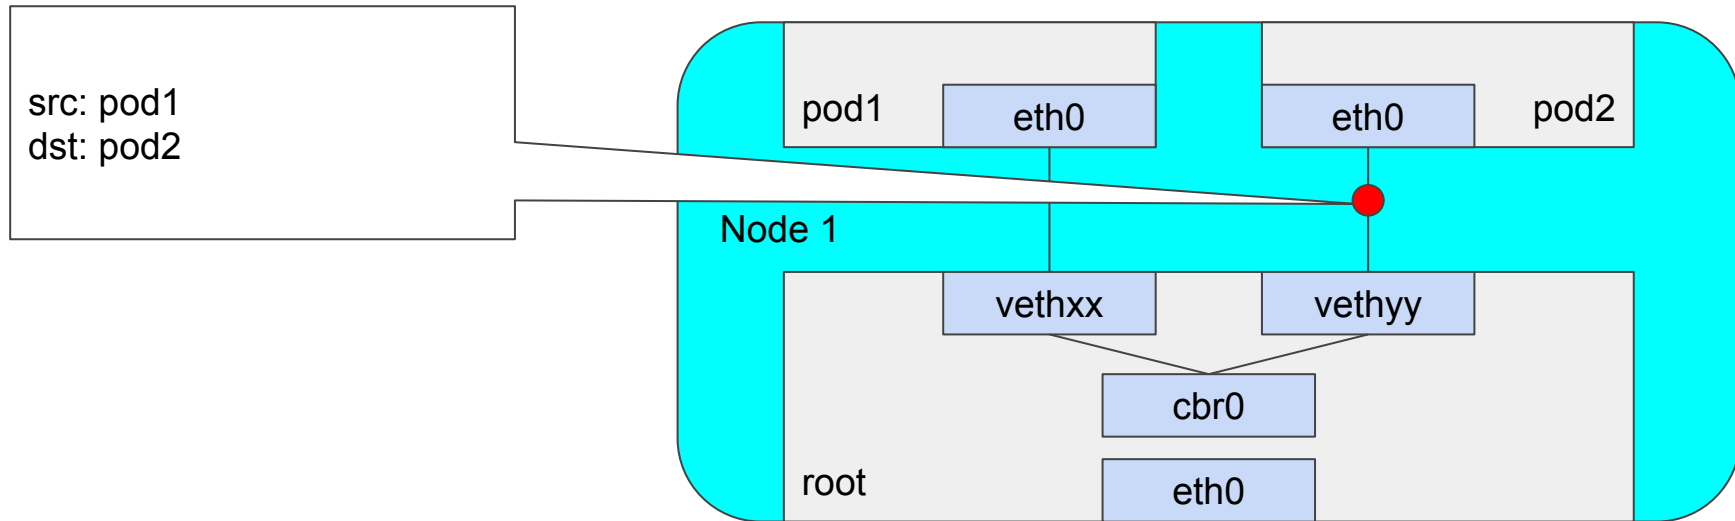
Pod to Pod - Intra Node



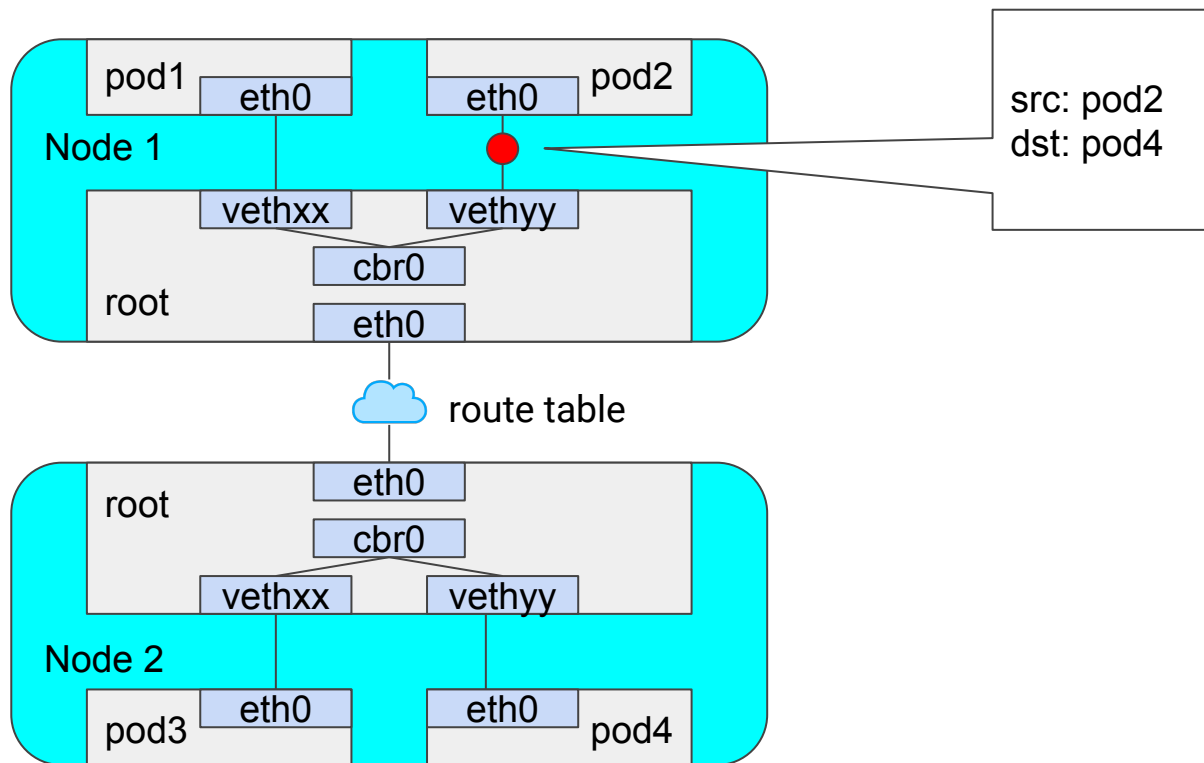
Pod to Pod - Intra Node




Pod to Pod - Intra Node




Pod to Pod - Inter Node



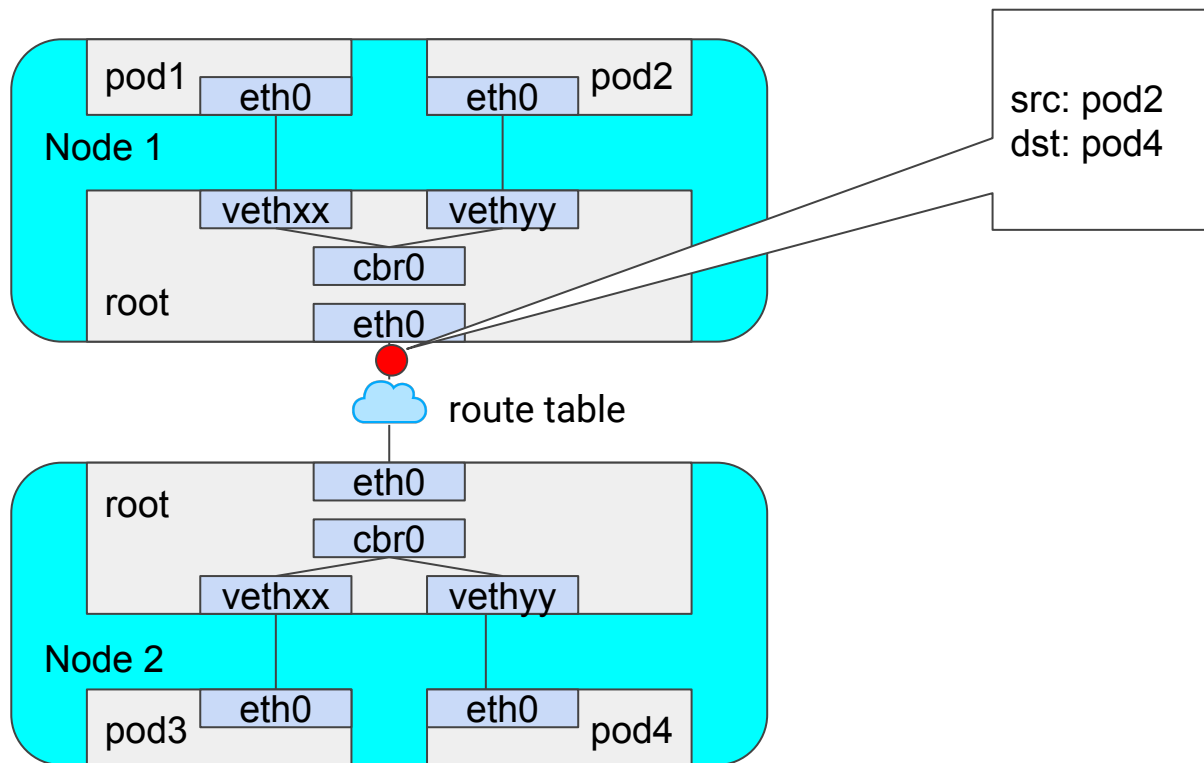
Pod to Pod - Inter Node



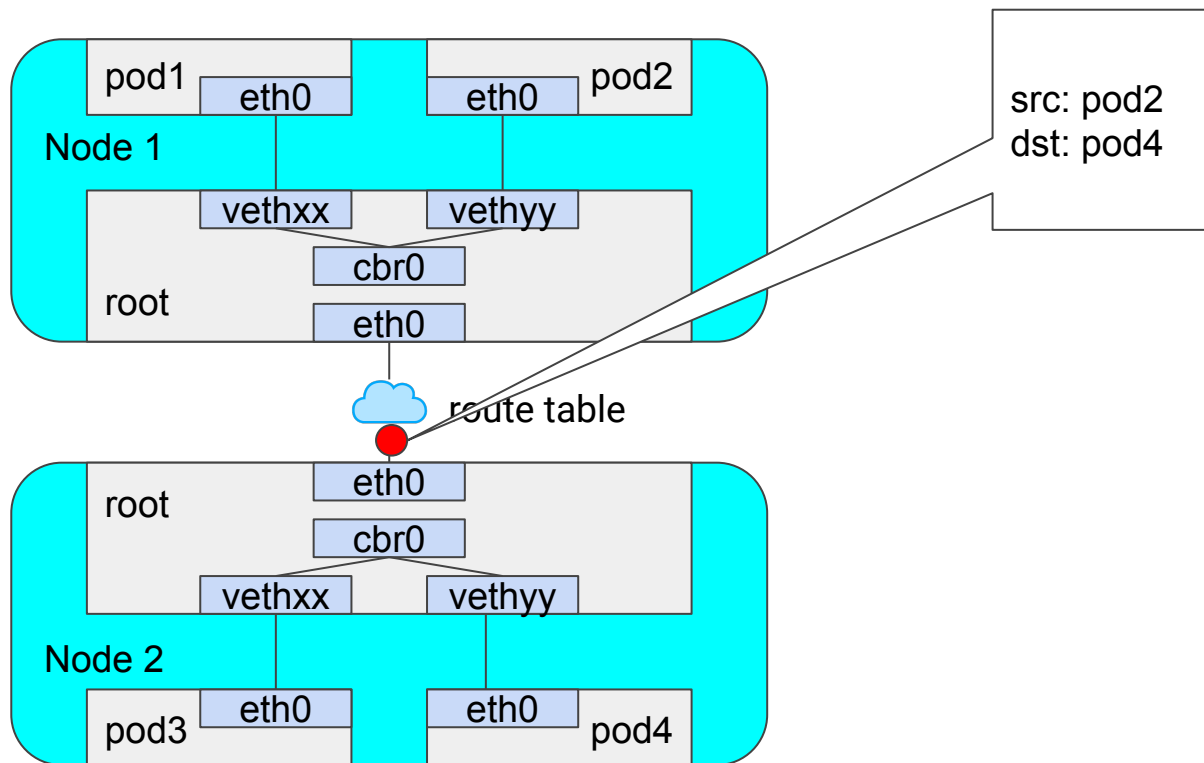
Pod to Pod - Inter Node




Pod to Pod - Inter Node




Pod to Pod - Inter Node




Pod to Pod - Inter Node



Pod to Pod - Inter Node



Pod to Pod - Inter Node




Basic is never enough sometimes

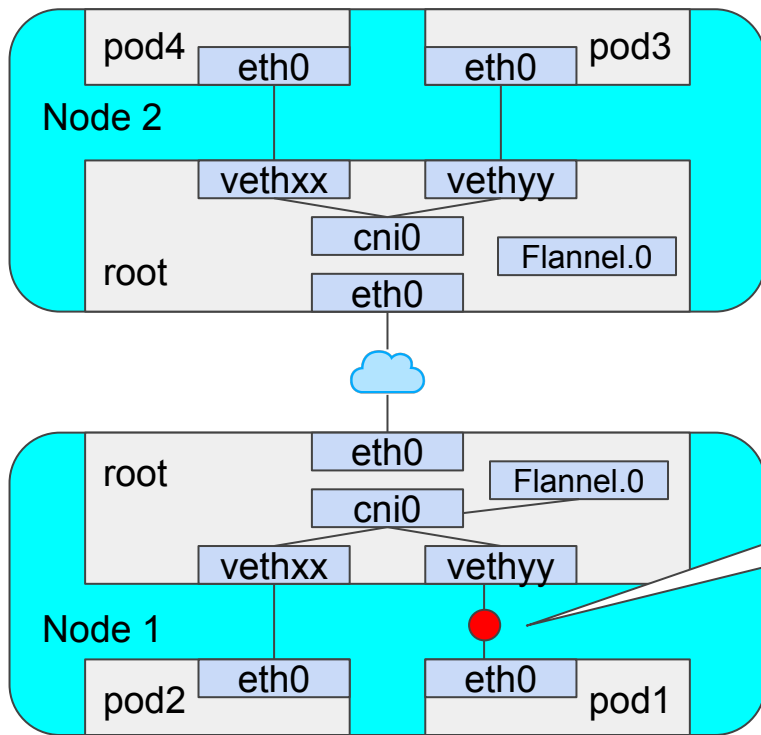
- depend on cloud provider route table capability (AWS VPC limit to 50 entry-> 50 node max per cluster)
- cluster cannot be setup in private vpc -> kubernetes only support single route tables

Container Network Interface(CNI) Plugin

- Flannel
- Calico
- Romana
- Weavenet
- Canal
- Cilium



Pod to Pod - Inter Node-Flannel(concept)



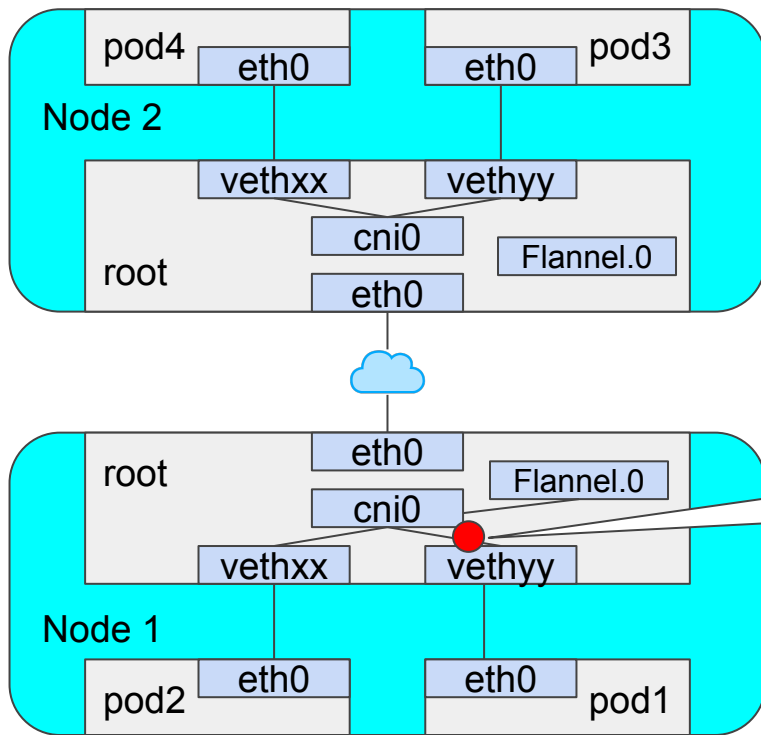
Overlay Network Solution(Encapsulate Packet to Packet)

solve the problem

- Not enough space IP
- Node Network cannot handle too much routing
- Additional management feature

src: pod1
dst: pod4

Pod to Pod - Inter Node-Flannel(concept)



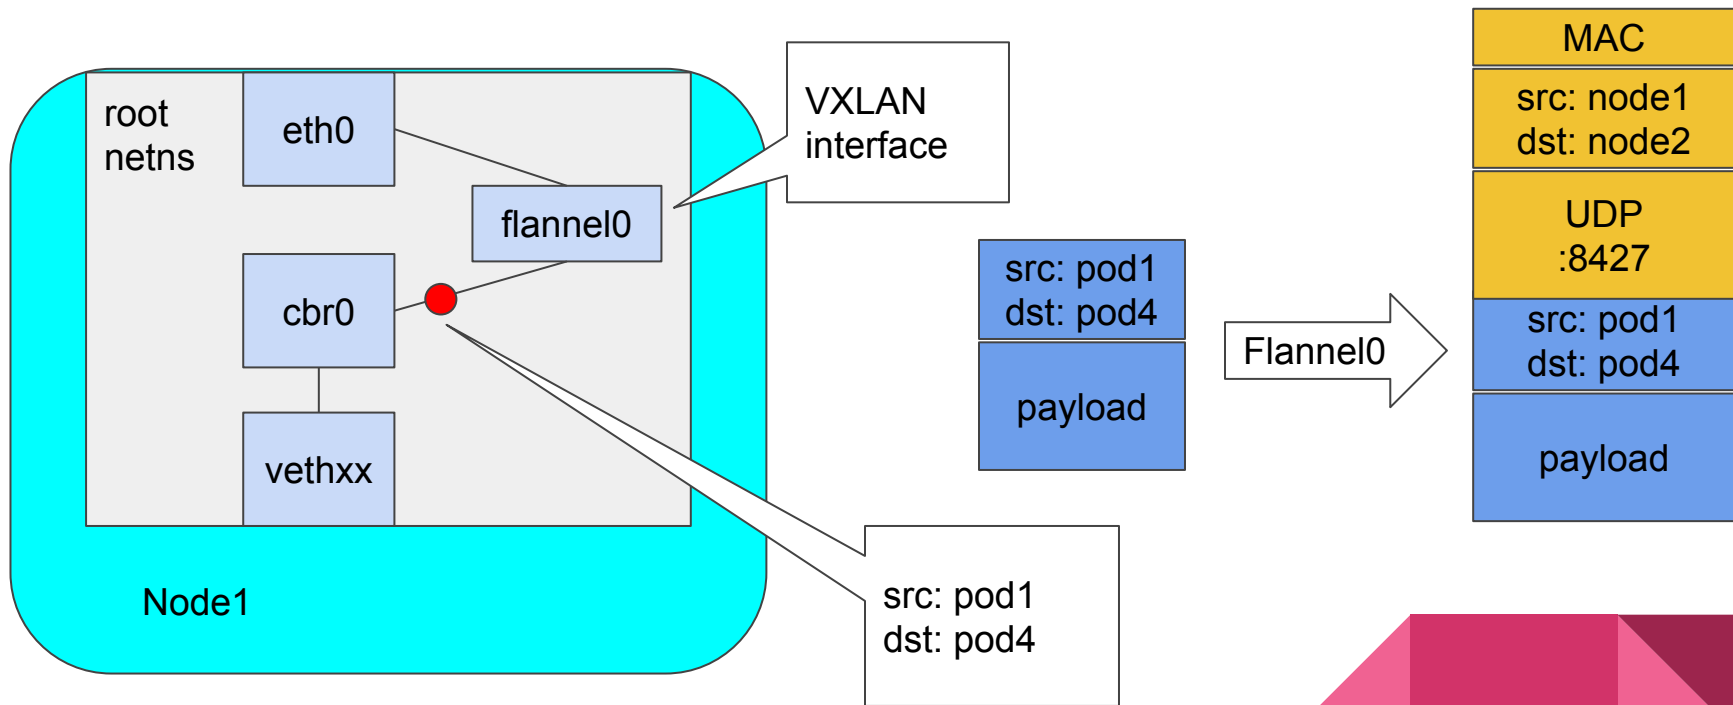
Overlay Network Solution(Encapsulate Packet to Packet)

solve the problem


- Not enough space IP
- Node Network cannot handle too much routing
- Additional management feature

src: pod1
dst: pod4

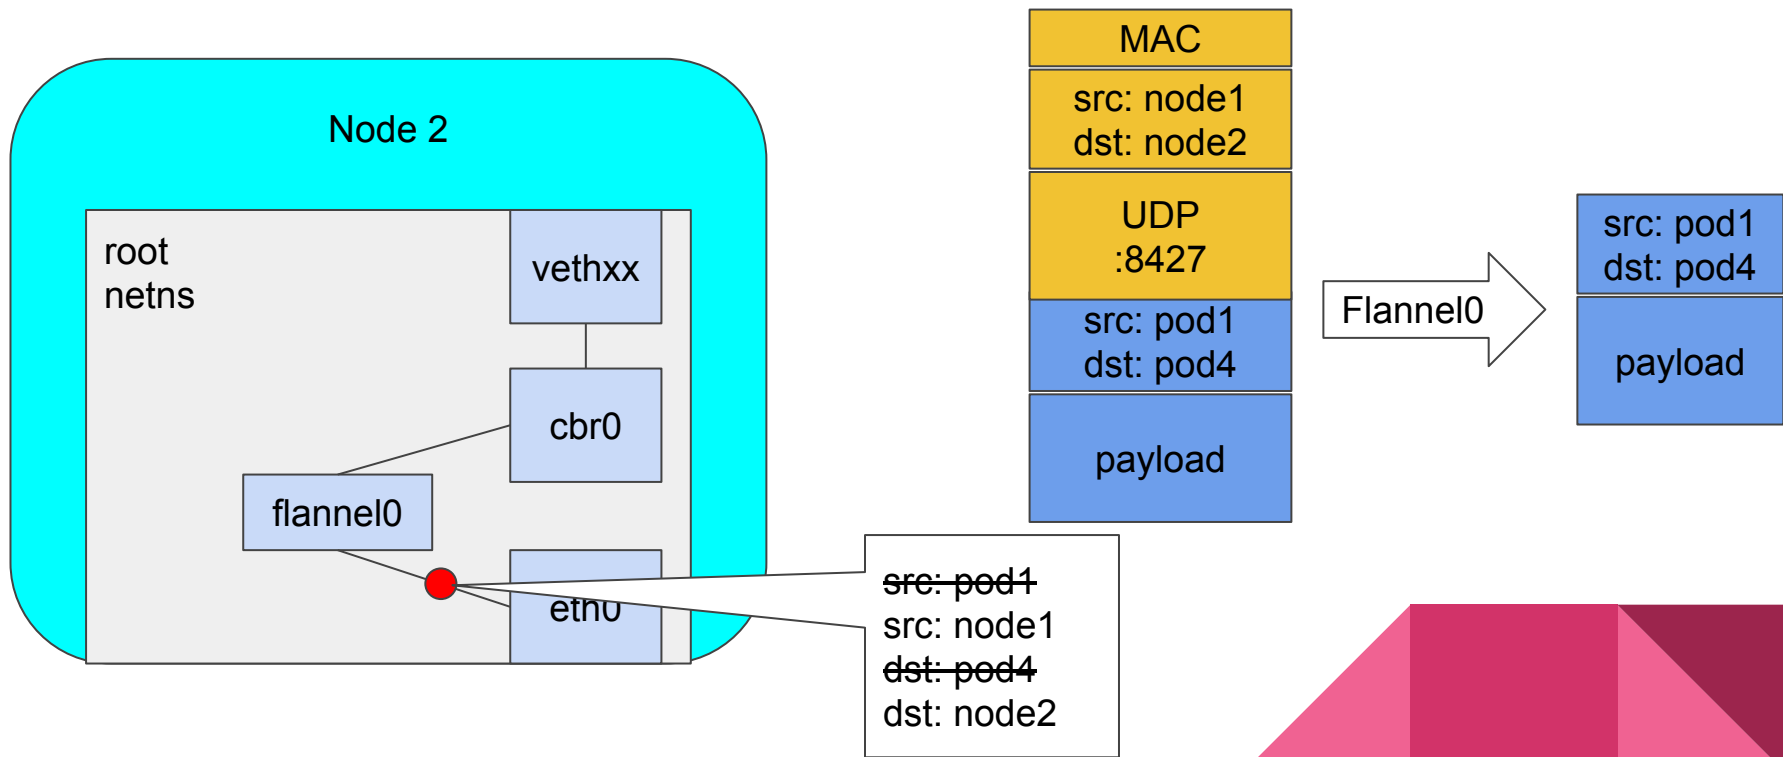
Pod to Pod - Inter Node-Flannel(concept)



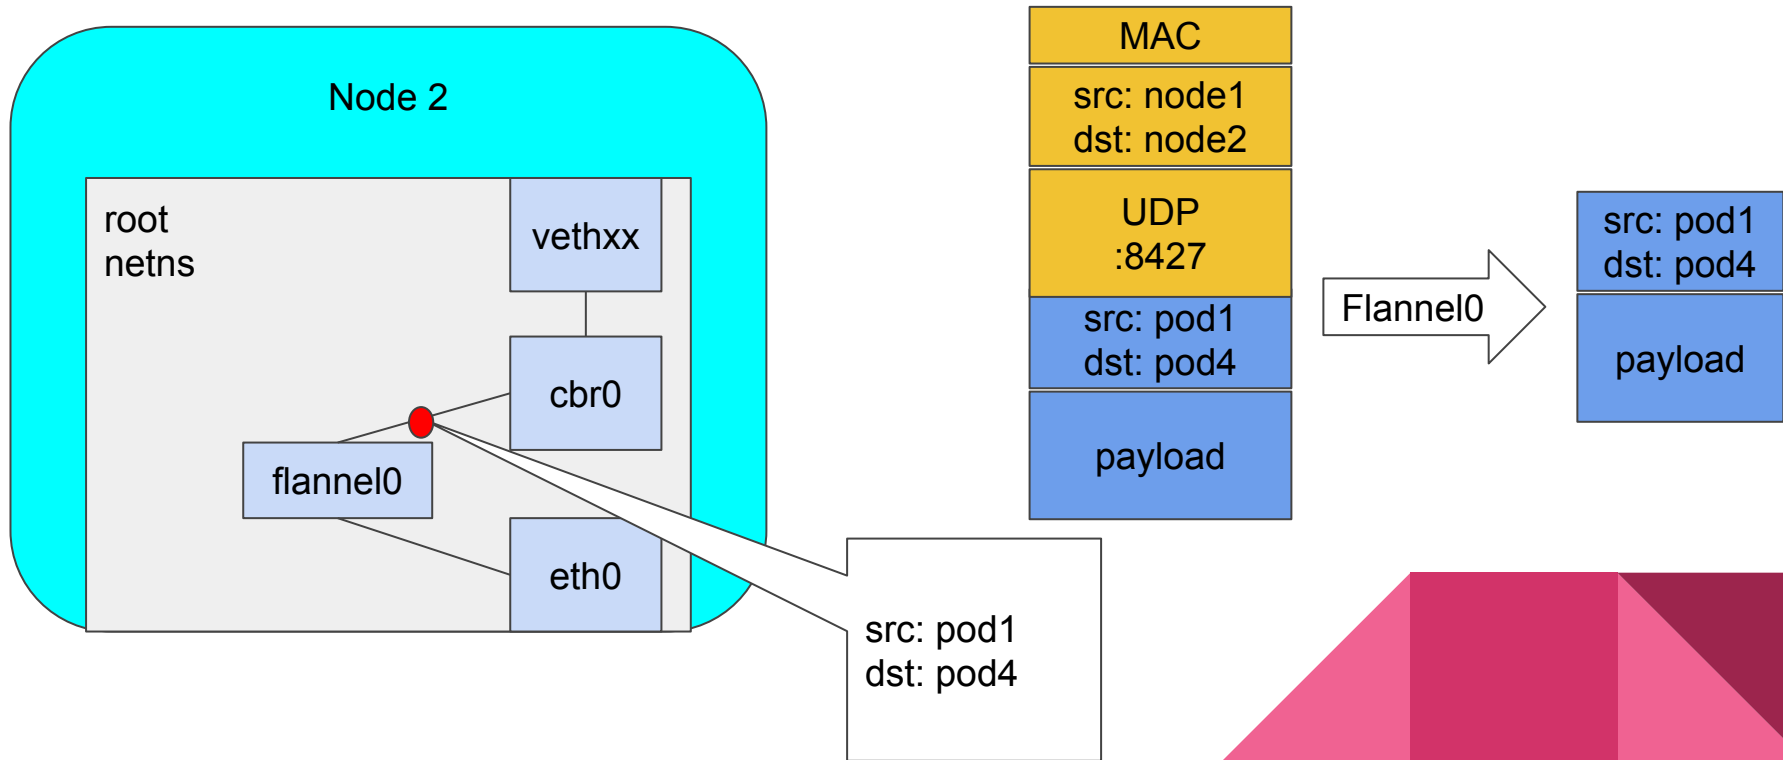
Pod to Pod - Inter Node-Flannel(concept)



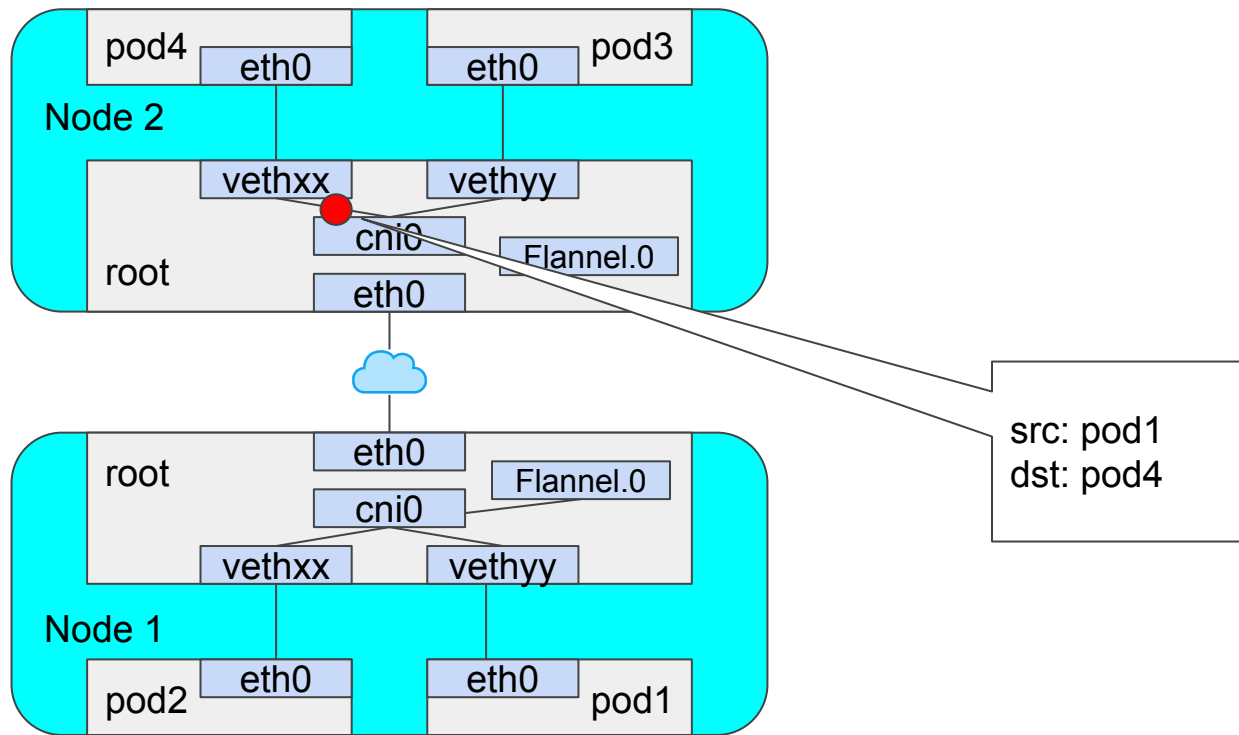
Pod to Pod - Inter Node-Flannel(concept)




Pod to Pod - Inter Node-Flannel(concept)



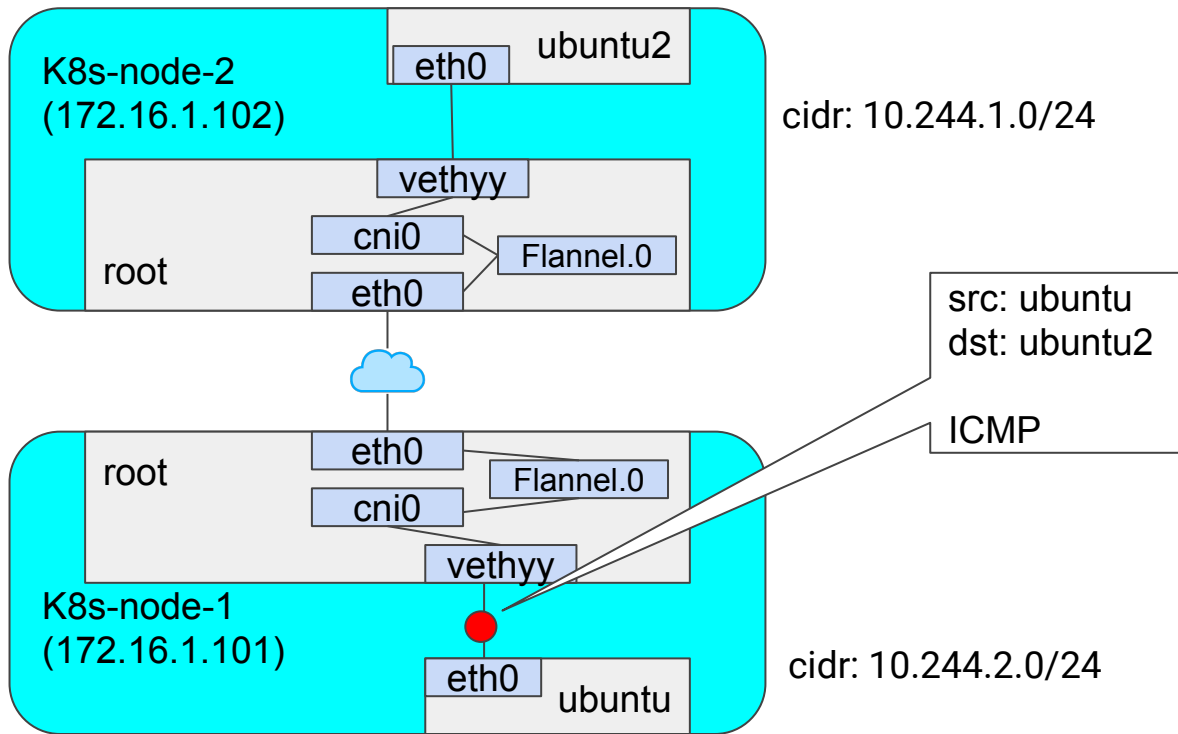
Pod to Pod - Inter Node-Flannel(concept)



Pod to Pod - Inter Node-Flannel(concept)



Pod to Pod - Inter Node-Flannel(demo)




Life is never flat...

rolling update
scale up/down
pods crash/hang
nodes reboot

pods is mortal. come and go, never get
ressurrected

we need some interface abstraction to handle
dynamics in pod life rotation






Services

Abstraction

- *load balancer pods*
- Stable Virtual IP
- Group of Endpoint
- Support Port Forwarding



Link Service to Pods(demo)

```
apiVersion:
extensions/v1beta1
kind: Deployment
metadata:
  name: deploy-nginx
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: proxy
    ----
```

```
apiVersion: v1
kind: Service
metadata:
  name: srv-nginx
spec:
  ports:
    - port: 80
      protocol: TCP
      targetPort: 80
  selector:
    app: proxy
    ----
```



```
apiVersion: v1
kind: Service
metadata:
  --
  name: srv-nginx
  namespace: default
  --
spec:
  clusterIP: 100.68.232.5
  ports:
    - port: 80
      protocol: TCP
      targetPort: 80
  selector:
    app: proxy
    sessionAffinity: None
  type: ClusterIP
```

Services Type

- **ClusterIp:**

assigned IP inside cluster and only accessible from inside only

- **NodePort**

service mapped into particular node port and node IP

- **LoadBalancer**

exposed to external traffic.

- **ExternalName**


without pods selector, access external resource or other namespace




Endpoints

- API Object that automatically created when deploy service
- list all pods below particular service


```
apiVersion: v1
kind: Endpoints
---
subsets:
- addresses:
  - ip: 10.244.3.1
    nodeName: k8s-node-1
    targetRef:
      -----
  - ip: 10.244.3.2
    nodeName: k8s-node-2
    targetRef:
      -----
```




10.244.3.1




10.244.3.2




10.244.1.3



10.244.1.2




10.244.2.3




10.244.2.4


Pod to Service



Pod to Service



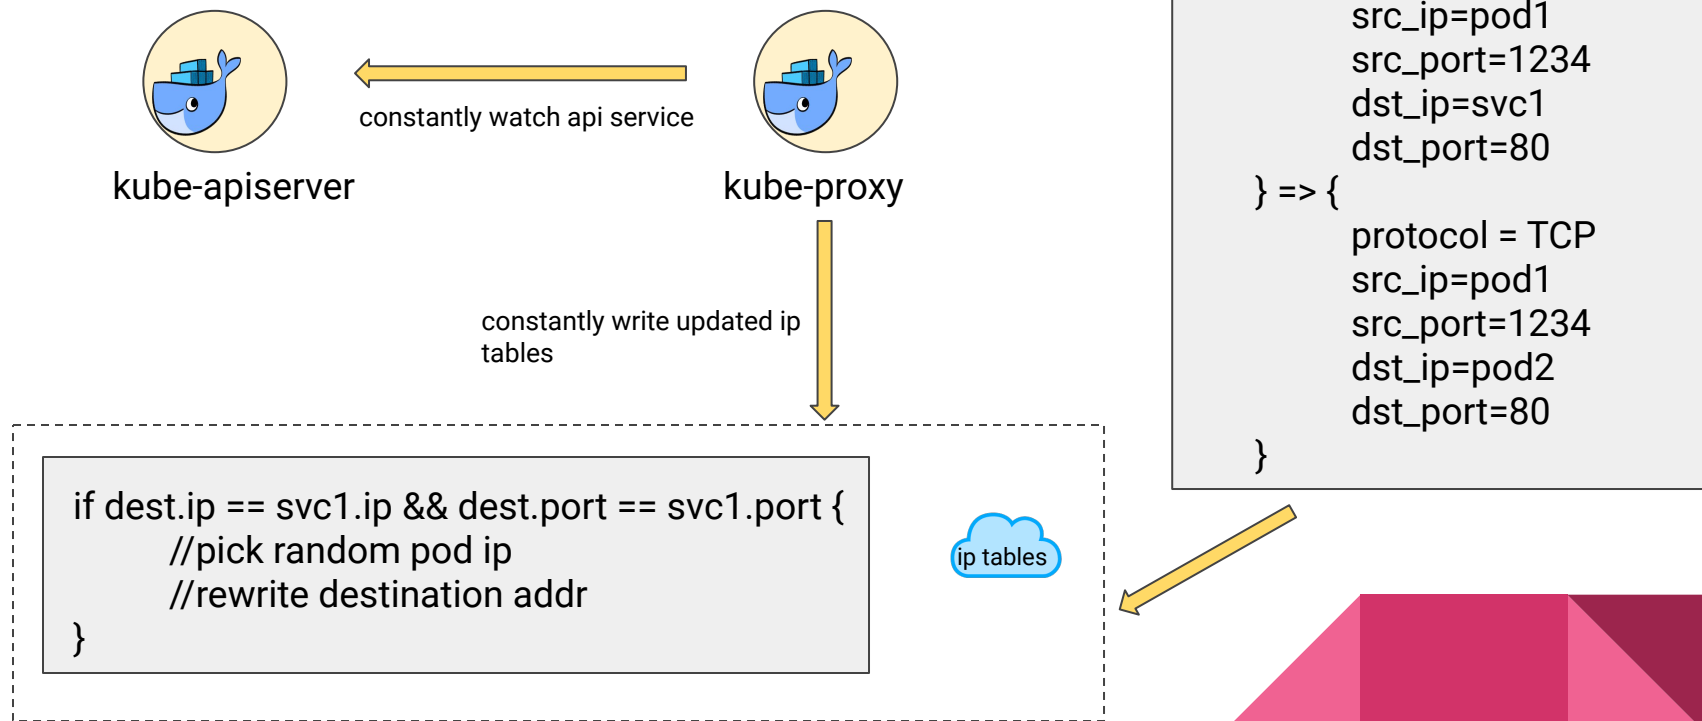
Pod to Service




src:pod1
~~dst:svc2~~
dst:pod2

DNAT Conntrack


ip tables & conntrack




Pod to Service




Pod to Service




Pod to Service



Pod to Service




Pod to Service

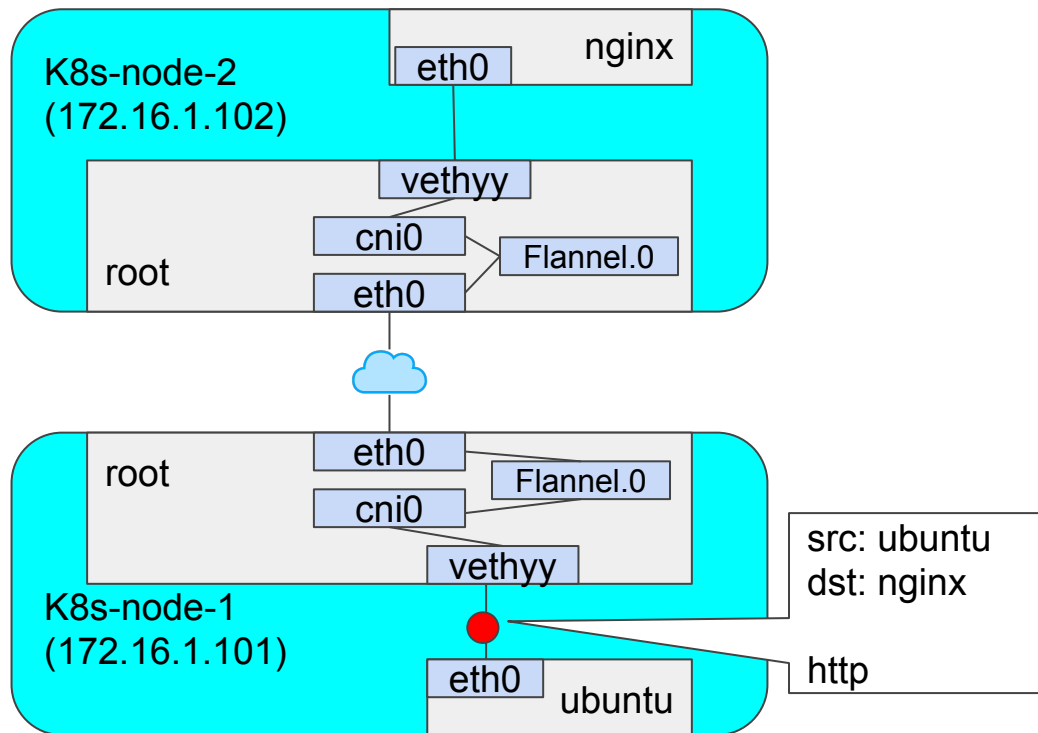


DNS

- say no to hardcode ip service. use friendly hostname, bonded with endpoint resource
- provide “A” & “SRV” records
- run in pods as kube-dns




Pod to Service(demo)




Cluster to External Communication


Send to External Traffic(Egress) - GCP Case



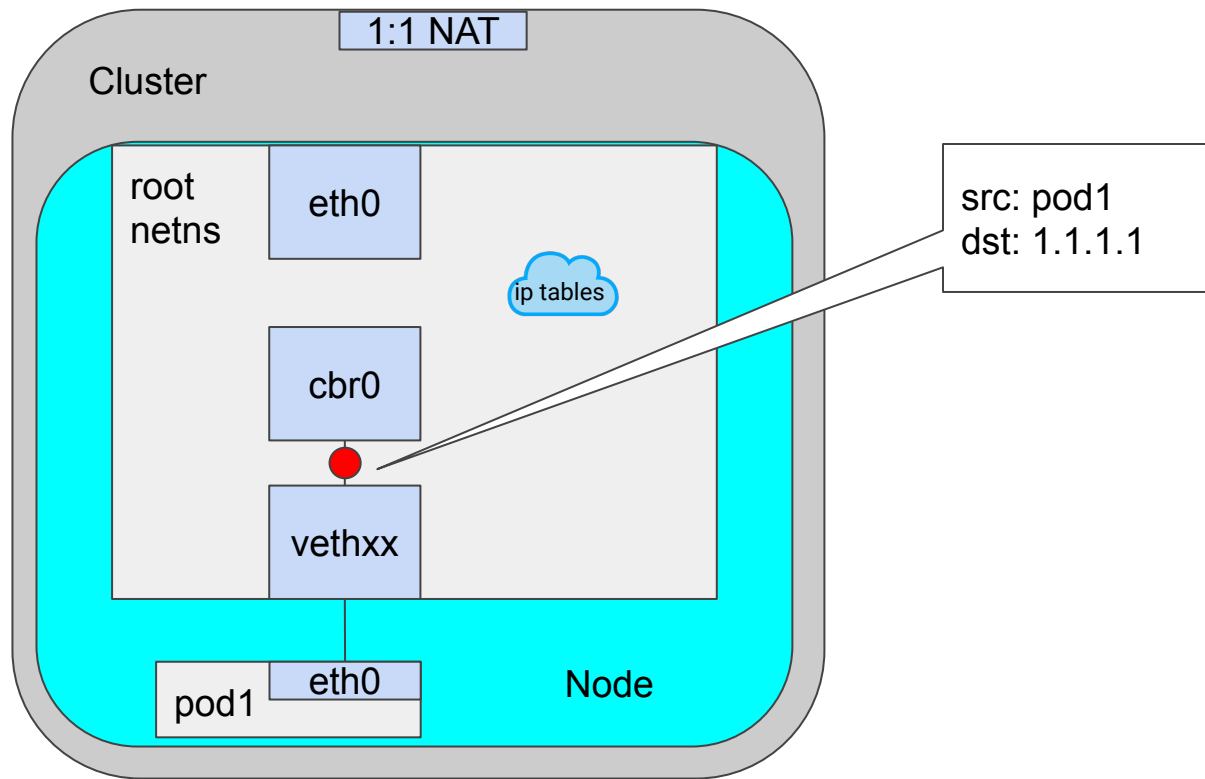
IP Private(inter-node) and maybe have public IP also




Life of packet: pod-to-internet




Life of packet: pod-to-internet



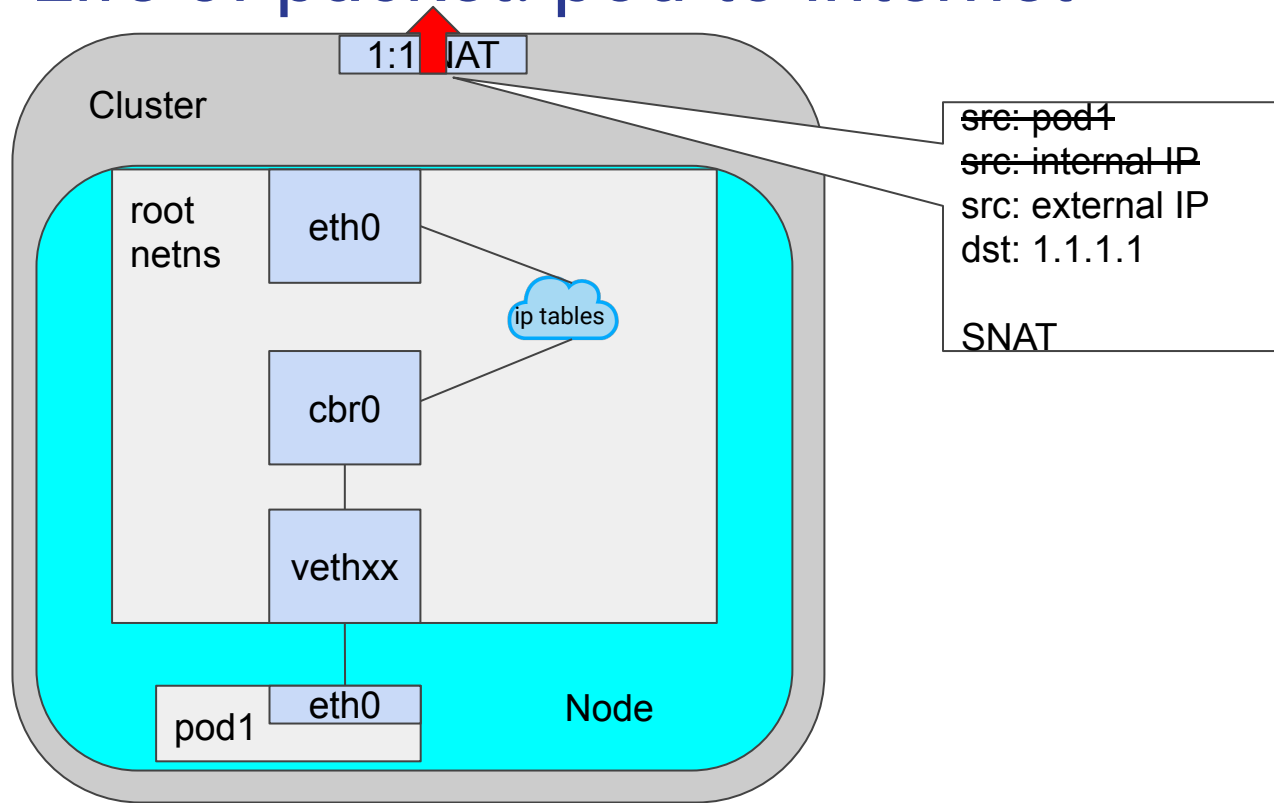
Life of packet: pod-to-internet



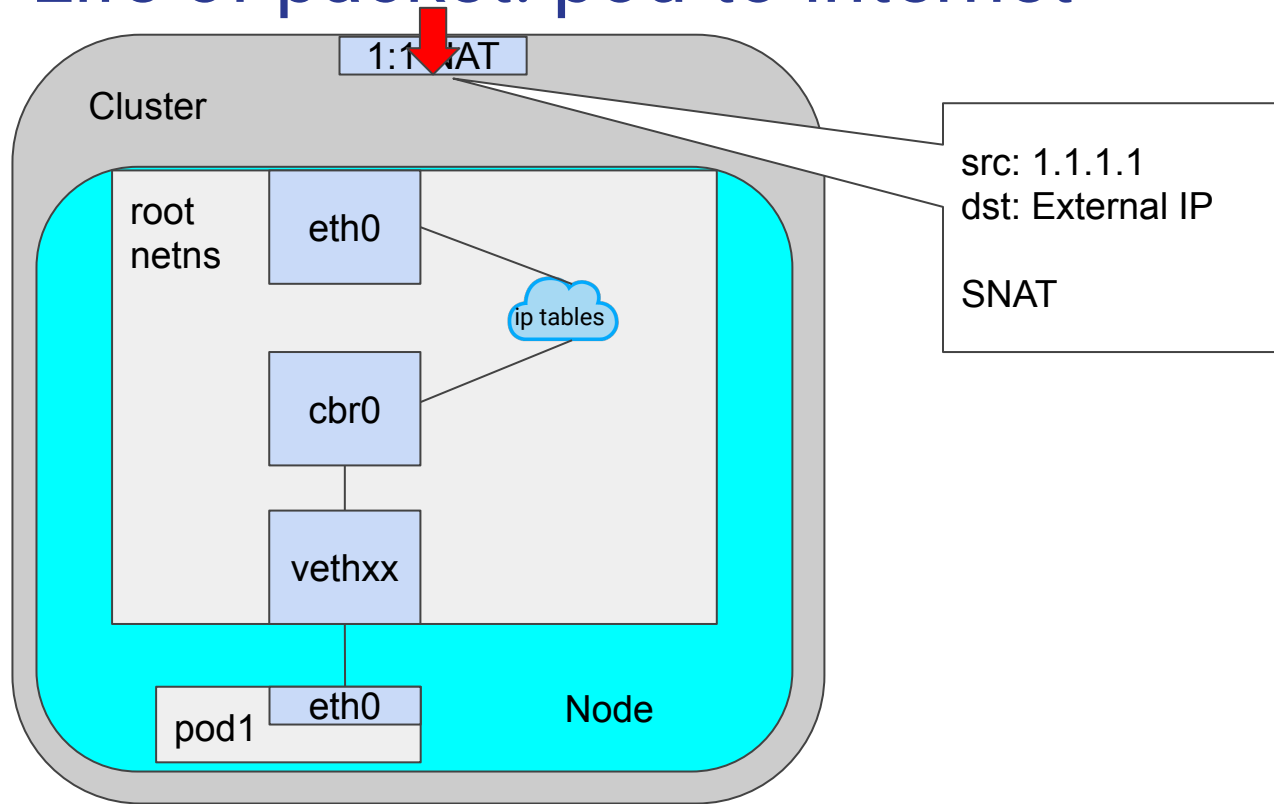
Life of packet: pod-to-internet




Life of packet: pod-to-internet



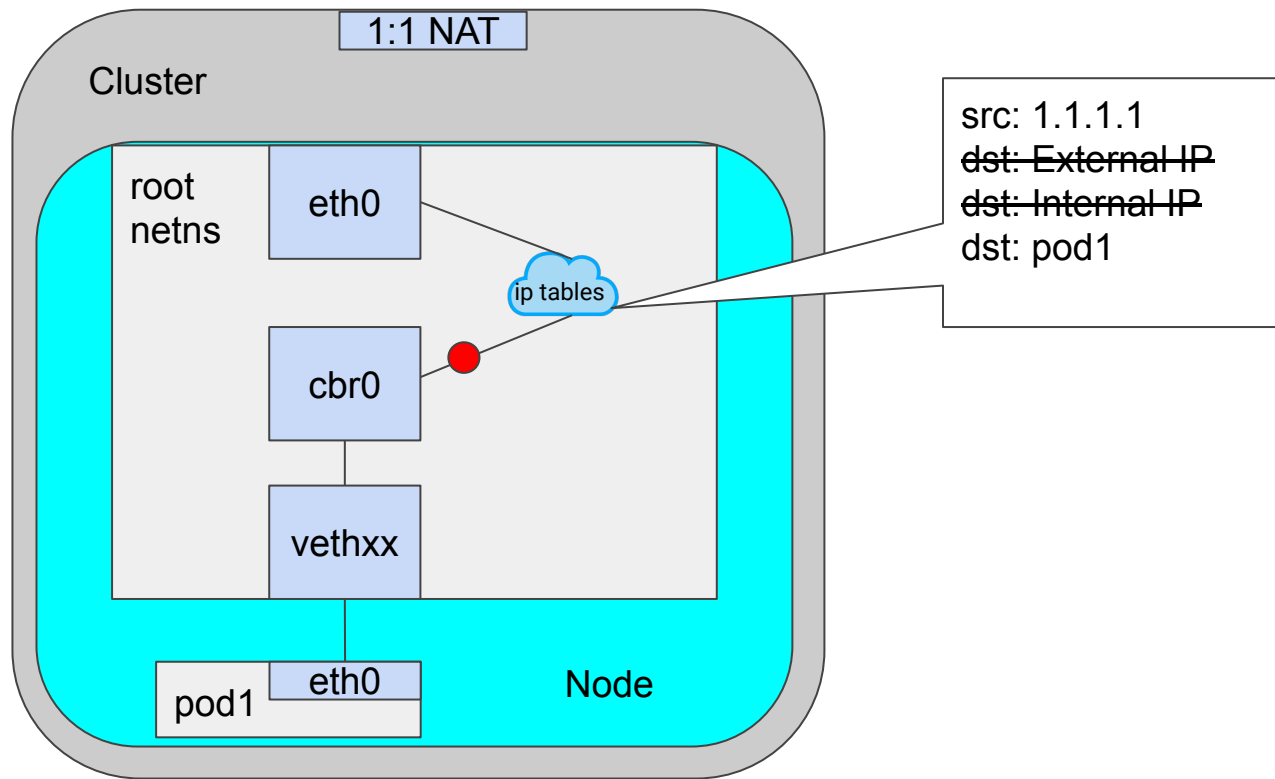
Life of packet: pod-to-internet




Life of packet: pod-to-internet



Life of packet: pod-to-internet



Life of packet: pod-to-internet



Receive from External Traffic(Ingress)


k8s support

L4: TCP


L7: HTTP/S (GCP)

mapped to


- Service type=LoadBalancer
- Ingress




Life of packet: external to service(L4 case)




Life of packet: external to service(L4 case)




Life of packet: external to service(L4 case)




Life of packet: external to service(L4 case)




Life of packet: external to service(L4 case)



Life of packet: external to service(L4 case)



Life of packet: external to service(L4 case)



Network Policy


Network Policy

like “firewall”, regulate inbound and outbound for pods

grouped by pods selector


default: allow all

policies can be stacked. order is important



Namespace

private area to manage
object(pods, service, policies)



Network Policy

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: example-network-policy
  namespace: default
spec:
```

```
  podSelector:
    matchLabels:
      role: db
```

```
  policyTypes:
    - Ingress
    - Egress
```

ingress:

```
- from:
  - podSelector:
      matchLabels:
        role: frontend
  ports:
    - protocol: TCP
      port: 6379
```

egress:


```
- to:
  - ipBlock:
      cidr: 10.0.0.0/24
  ports:
    - protocol: TCP
      port: 597
```

- need to specify
- subject pods
 - rule ingress/egress
 - open port



Custom Network Policy

- CNI that support additional feature in Network Policy
 - calico
 - romana
 - weavenet
 - openshift
 -



Thats All...

Pahamify is Hiring..

pahamify.com/career

or

Contact:

career@pahamify.com



