

Basic of Stateful and Storage in Kubernetes

Kevin Setiawan Tanzil

Software Engineer @Alterra

1 August 2019





Agenda

How do these words mean and how they work together

Ephemeral

Remote File

Out-of-Tree

Stateful

Block

Dynamic Provisioning

Persistent Volume Claims

CSI

In Tree

Persistent Volume

Stateless

Volume

Storage Classes

Local

Object

Driver


Plugin



Container ?

What is the problem ?


1. Container are ephemeral
2. Container can't share data from each other





The Solution is **Kubernetes Volume**

1. Accessible by all container in pods
2. Lifetime of volume may same with pod or longer





Kubernetes Volume Plugins






Kubernetes Volume Plugins

What is that ?

1. A way to reference **block device** or **mounted file system**
2. Volume plugins specify :
 - a. How volume is set up inside a pod
 - b. Medium that back it up





Kubernetes Volume Plugins

Currently Supported

1. Remote Storage

- a. GCE Persistent Disk
- b. AWS Elastic Block Store
- c. Azure File Storage
- d. NFS
- e. Ceph
- f. vSphere
- g. iSCSI
- h. Etc.

2. Ephemeral Storage

- a. EmptyDir

3. Local Persistent Volume

4. Out-of-tree

- a. CSI (Container Storage Interface)
- b. Flex


5. Others

- a. Host path

Kubernetes Volume Plugins

Ephemeral Storage

1. Basically is a temp scratch file system on host machine
2. Life is same with pod lifetime, pod destroy volume also destroy
3. Can only be reference in-line on pod cant via PV/PVC
4. Volume Plugins : EmptyDir





Kubernetes Volume Plugins

Ephemeral Storage

1. Basically is a temp scratch file system on host machine
2. Life is same with pod lifetime, pod destroy volume also destroy
3. Can only be reference in-line on pod cant via PV/PVC
4. Volume Plugins : EmptyDir

```
apiVersion: v1
kind: pod
metadata:
  name: test-pod
spec:
  containers:
    - image: example/container1
      name: container1
      volumeMounts:
        - mountPath: /shared
          name: shared-emptydir-space
    - image: example/container2
      name: container2
      volumeMounts:
        - mountPath: /shared
          name: shared-emptydir-space
  volumes:
    - name: shared-emptydir-space
      emptyDir: {}
```



Kubernetes Volume Plugins


Remote Storage

1. Lifetime data persists beyond of any pod lifecycle
2. Reference in-line or via PV/PVC
3. Example of Remote Storage :
 - a. GCE Persistent Disk
 - b. AWS Elastic Block Store
 - c. Azure Data Disk
 - d. Ceph
 - e. And More....

Kubernetes Volume Plugins

Remote Storage

1. Kubernetes Will Attach
a. Attach Volume
b. Mount Volume



pv1

pypod

ta-remote

stentDisk:

: panda-disk

: ext4

:

rc.io/google_container/busybox

leepycontainer

p

0"

unts:

- mountPath: /data

name: data

readOnly: false



Persistent Volumes & Persistent Volume Claims





Persistent Volume

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: mypv1
spec:
  accessModes:
    - ReadWriteOnce
  capacity:
    storage: 10Gi
  persistentVolumeReclaimPolicy: Retain
  gcePersistentDisk:
    fsType: ext4
    pdName: panda-disk-small
```

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: mypv2
spec:
  accessModes:
    - ReadWriteOnce
  capacity:
    storage: 100Gi
  persistentVolumeReclaimPolicy: Retain
  gcePersistentDisk:
    fsType: ext4
    pdName: panda-disk-big
```



Persistent Volume Claims

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mypvc
  namespace: testns
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 100Gi
```



PV to PVC Binding

```
kevinsetiawantanzil2@cloudshell:~/kubetalk-test (kubetalk)$ kubectl create -f small-pv.yaml
persistentvolume/mypv1 created
kevinsetiawantanzil2@cloudshell:~/kubetalk-test (kubetalk)$ kubectl create -f big-pv.yaml
persistentvolume/mypv2 created
kevinsetiawantanzil2@cloudshell:~/kubetalk-test (kubetalk)$ kubectl get pv
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	REASON	AGE
mypv1	10Gi	RWO	Retain	Available		manual		9s
mypv2	100Gi	RWO	Retain	Available		manual		4s

```
kevinsetiawantanzil2@cloudshell:~/kubetalk-test (kubetalk)$ kubectl create -f mypvc.yaml
persistentvolumeclaim/mypvc created
kevinsetiawantanzil2@cloudshell:~/kubetalk-test (kubetalk)$ kubectl get pv
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	REASON	AGE
mypv1	10Gi	RWO	Retain	Available		manual		26s
mypv2	100Gi	RWO	Retain	Bound	default/mypvc	manual		21s



Kubernetes Volume Plugins

Remote Storage

1. Volume Reference via PVC
2. POD YAML is Portable across cluster

```
apiVersion: v1
kind: pod
metadata:
  name: sleepypod
spec:
volumes:
  name: data-remote
  gcePersistentDisk:
    pdName: panda-disk
    fsType: ext4
  volume:
    - name: data
      persistentVolumeClaim:
        claimName: mypvc
containers:
  - image: grc.io/google_container/busybox
    name: sleepycontainer
    command:
      - sleep
      - "6000"
  volumeMounts:
    - mountPath: /data
      name: data
      readOnly: false
```




Dynamic Provisioning & Storage Classes





Dynamic Provisioning

1. Administrator admin pre-provisioning PVs is painful and wasteful
2. Create new volume on demand
3. Remove administrator for pre-provisioning PVs



Dynamic Provisioning Storage Classes

1. Dynamic Provisioning Enabled by making StorageClass
2. StorageClass define parameter that needed
3. StorageClass parameter is opaque

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
```

```
  name: slow
  provisioner:
    kubernetes.io/gce-pd
  parameters:
```

```
    type: pd-standard
```

```
--
```

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
```

```
  name: fast
  provisioner:
    kubernetes.io/gce-pd
  parameters:
```

```
    type: pd-ssd
```



Dynamic Provisioning

Persistent Volume Claims

1. Select storageClass in PVC will trigger Dynamic Provisioning to create PV

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mypvc
  Namespace: testns
spec:
  accessModes:
  - ReadWriteOnce
  resource:
    request: 100Gi
  storageClassName: fast
```



Dynamic Provisioning

Remote Storage

```
apiVersion: v1
kind: pod
metadata:
  name: sleepypod
volume:
- name: data
  persistentVolumeClaim:
    claimName: mypvc
containers:
- image: grc.io/google_container/busybox
  name: sleepycontainer
  command:
    - sleep
    - "6000"
  volumeMounts:
    - mountPath: /data
      name: data
      readOnly: false
```



Dynamic Provisioning

Default Storage Classes

1. Enable Dynamic Provisioning even StorageClass not specific
2. Pre-install Storage Class

```
kind: StorageClass
apiversion: storage.k8s.io/v1
metadata:
  name: slow
  annotations:
```

```
storageclass.beta.kubernetes.io/is-default-class:"true"
provisioner: kubernetes.io/gce-pd
parameters:
  type: pd-standard
```

```
--
```

```
kind: StorageClass
apiversion: storage.k8s.io/v1
metadata:
  name: fast
provisioner: kubernetes.io/gce-pd
parameters:
  type: pd-ssd
```



In-Tree Volume Plugins

What is this ?

1. A list of Supported Volume plugins by Kubernetes called In Tree Volume Plugins
2. In-Tree Volume plugins are awesome
 - a. Can do dynamic provisioning
 - b. Automate provisioning
 - c. Automate mounting

Cons :

1. Kubernetes dev should maintain the volume plugins code
2. A little bug in volume plugin can make critical impact to kubernetes component
3. Storage Vendor should keep update with kubernetes release.
4. Vendor force to be open source



Out-of-Tree Volume Plugins

CSI (Container Storage Interface)

1. The idea is makes the Volume Plugins Truly extensible
2. Plugin may be containers
3. Storage Vendor can develop Driver base on CSI





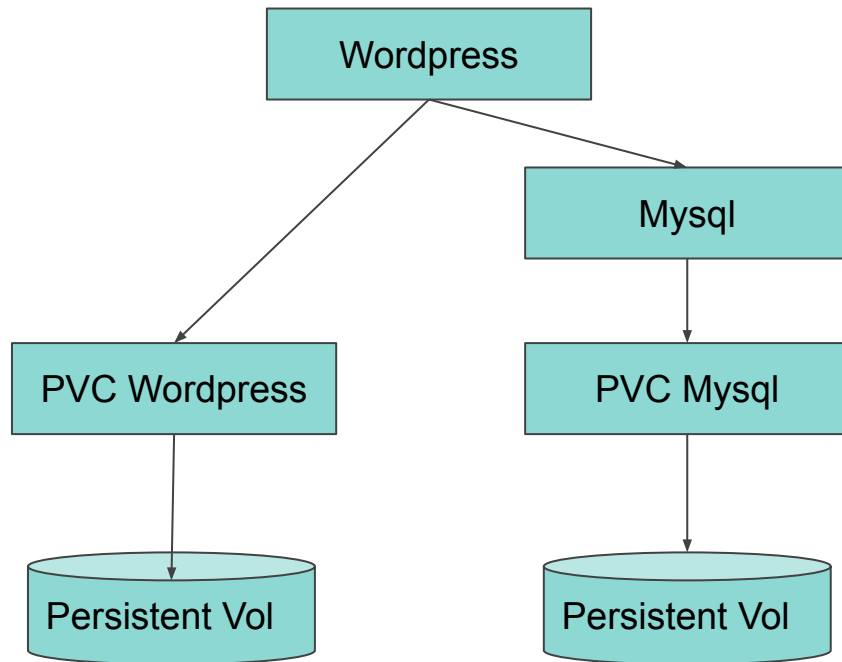
DEMO





What we're gonna to make

Apps Diagram



Thank you

Github : github.com/tanzilgr2288

Medium : medium.com/@kevinsetiawantanzil2